

A reliable Turing machine

Ikir Çapuni

Peter Gács

May 30, 2019

Abstract

The title says it.

Contents

1	Introduction	2
1.1	To be written	2
1.2	Turing machines	2
1.3	Codes, and the result	4
1.4	A shortcut solution	5
2	Overview of the construction	5
2.1	Isolated bursts of faults	5
2.2	Hierarchical construction	7
2.3	From combinatorial to probabilistic noise	8
2.4	Difficulties	9
3	Notation	12
4	Specifying a Turing machine	13
4.1	Universal Turing machine	13
4.2	Rule language	15
4.3	Fields	16
5	Exploiting structure in the noise	19

5.1	Sparsity	19
5.2	Error-correcting code	21
6	The model	22
6.1	Generalized Turing machine	22
6.2	Simulation	26
6.3	Hierarchical codes	27
6.4	Trajectories	28
7	Simulation structure	30
7.1	Head movement	31
7.2	Computation phase	35
7.3	Forced self-simulation	37
7.4	Transfer phase	39
8	Health	40
9	Healing and rebuilding	44
9.1	Annotation, admissibility, stitching	44
9.2	The healing procedure	47
9.3	Rebuilding	49
10	Scale-up	52
10.1	Super-health	52
10.2	Annotated history	53
10.3	The simulation codes	54
11	Isolated bursts	55
12	Cleaning	58
12.1	Cleaning the current level	58
12.2	Escape	60
12.3	Pass cleaning	65
12.4	Spill bound	68
12.5	Attack cleaning	70
13	After a large burst	72
13.1	Spill bound	72

1 Introduction

1.1 To be written

1.2 Turing machines

Our contribution could use any one of the standard definitions of a Turing machine, but it will be convenient to modify it. Let us recall that a one-tape Turing

machine is defined by a finite set Γ of *internal states*, a finite alphabet Σ of *tape symbols*, a transition function

$$\delta: \Sigma \times \Gamma \rightarrow \Sigma \times \Gamma \times \{-1, 1\},$$

and possibly some distinguished states and tape symbols. Any array $A \in \Sigma^{\mathbb{Z}}$ is called the *tape configuration*, or *tape content*. At any time, the head is at some integer position h , and is observing the tape symbol $A(h)$. The meaning of $\delta(a, q) = (a', q', d)$ is that if $A(h) = a$ and the state is q then the $A(h)$ will be rewritten as a' and h will change to $h + d$.

We will use a model that is slightly different, but has clearly the same expressing power. There is no set of internal states, but the head observes and modifies *two* neighboring tape cells at a time. Thus, a Turing machine is defined as a pair

$$(\Sigma, \tau).$$

The tape alphabet Σ contains at least the distinguished symbols $\sqcup, 0, 1$ where \sqcup is called the *blank symbol*. We have a *transition function*

$$\tau: \Sigma^2 \rightarrow \Sigma^2 \times \{-1, 1\}.$$

The tape is blank at all but finitely many positions. A *configuration* is a pair

$$(A, h),$$

$h \in \mathbb{Z}$ is the current *head position*, or *observed cell*, or *current cell*, and $A \in \Sigma^{\mathbb{Z}}$ is the *tape content*, or *tape configuration*: at position p , the tape contains the symbol $A(p)$. If $\xi = (A, h)$ is a configuration then we will write

$$\xi.\text{tape} = A, \quad \xi.\text{pos} = h. \tag{1.1}$$

Though the tape alphabet may contain non-binary symbols, we will restrict input and output to binary.

The transition function τ tells us how to compute the next configuration from the present one. When the head observes the pair of tape cells with content $\mathbf{a} = (a_0, a_1)$ at positions $h, h + 1$ then denoting

$$(\mathbf{a}', j) = \tau(\mathbf{a}),$$

will change tape content at positions $h, h + 1$ to a'_0, a'_1 , and move the head to tape position to $h + j$.

Definition 1.1 (Fault) A *fault* occurs at time t if the output (\mathbf{a}', j) of the transition function at this time is replaced with some other value (which is then used to compute the next configuration). \dashv

1.3 Codes, and the result

For fault-tolerant computation, some redundant coding of the information is needed.

Definition 1.2 (Codes) Let Σ_1, Σ_2 be two finite alphabets. A *block code* is given by a positive integer Q —called the *block size*—and a pair of functions

$$\psi_* : \Sigma_2 \rightarrow \Sigma_1^Q, \quad \psi^* : \Sigma_1^Q \rightarrow \Sigma_2$$

with the property $\psi^*(\psi_*(x)) = x$. It is extended to strings by encoding each letter individually: $\psi_*(x_1, \dots, x_n) = \psi_*(x_1) \cdots \psi_*(x_n)$. \lrcorner

For ease of spelling out a result, we consider only computations whose outcome is a single symbol, at tape position 1. The machine will be declared *halted* if it writes the blank symbol at tape position 0.

With our machine we want to perform some computation, and want the result to be correct with large probability even if faults can occur with some small probability independently at each step. For this, the input will be encoded by some error-correcting code, to defend against the possibility of losing information from it even at the first reading. It is natural to allow the redundancy of the code to depend on the size of the input: in our case, it will depend logarithmically.

We will generally view a tape symbol as a tuple consisting of several *fields*. The output of computation should be possible read from one of these fields, which for symbol a will be called $a.output$.

Here is one formulation of the main result, about the ability to compute a function with values in $\{0, 1\}$ in a fault-tolerant way. The role of “halting” is played by the act of writing the blank symbol.

Theorem 1 *There is a Turing machine M_1 with a function $a \mapsto a.output$ defined on its alphabet, such that for any Turing machine G with alphabet Σ there are $0 \leq \varepsilon < 1$ and $\alpha_1, \alpha_2 > 0$ with the following property.*

For each input length $n = |x|$ a block code (φ_, φ^*) of block size $Q = O((\log n)^{\alpha_1})$ can be constructed such that the following holds.*

Let M_1 start its work from the initial configuration $\xi = (\varphi_(x), 0)$. Assume further that during its operation, faults occur independently at random with probabilities $\leq \varepsilon$.*

Suppose at time t the machine G writes the blank symbol at position 0, and value y at position 1. Then denoting by $\eta(u)$ the configuration of M_1 at time

u , at any time t' after

$$t \cdot (\log t)^{\alpha_2},$$

we have $\eta(t').\text{tape}(1).\text{output} = y$ with probability at least $1 - O(\epsilon)$.

We emphasize that the actual code φ of the construction will depend on n only in a simple way.

1.4 A shortcut solution

A fault-tolerant one-dimensional cellular automaton is constructed in [4]. If our Turing machine could just simulate such an automaton, it would become fault-tolerant. This can indeed almost be done provided that the size of the computation is known in advance. The cellular automaton can be made finite, and we could define a “kind of” Turing machine with a *circular tape* simulating it. But this solution requires input size-dependent hardware.

It seems difficult to define a fault-tolerant sweeping behavior on a regular Turing machine needed to simulate cellular automaton, without recreating an entire hierarchy—as we are doing here.

2 Overview of the construction

A Turing machine that simulates “reliably” any other Turing machine even when it is subjected to isolated bursts of faults of constant size, is given in [1]. By *reliably* we mean that the simulated computation can be decoded from the history of the simulating machine despite occasional damages. We will use some of the ideas of [1], but we will need to add new ones in order to maintain a hierarchy.

2.1 Isolated bursts of faults

Let us give a brief overview of a machine M_1 that can withstand isolated bursts of faults, as most of its construction will be reused in the probabilistic setting.

We break up the task of error correction into several problems to be solved. The solution of one problem gives rise to another one, but the process converges.

Redundant information The tape information of the simulated Turing machine will be stored in a redundant form, more precisely in the form of a block code.

Redundant processing The block code will be decoded, the retrieved information will be processed, and the result recorded. To limit the propagation of faults,

the tape will be split into tracks that can be handled separately, and the major processing steps will be carried out three times within one work period.

Local repair All the above process must resist a local burst of faults. For this, it is organized into a rigid, locally checkable structure with the help of local addresses, and some other tools like sweeps and short switchbacks (zigzags). The local healing procedure is the most complex part of the construction.

Disturbed local repair A careful organization of the healing procedure makes sure that even if a new burst interrupts it (or jumps into its middle), one or two new invocations of it will finish the job.

Here is some more detail. Each tape cell of the simulated machine M_2 will be represented by a block of some size Q called a *colony*, of the simulating machine M_1 . Each step of M_2 will be simulated by a computation of M_1 called a *work period*. During this time, the head of M_1 makes a number of sweeps over the current colony and its right neighbor colony, decodes the represented cell symbols, then computes and encodes the new symbols, and finally moves the head to the new position of the head of M_2 .

In order to protect information from the propagation of errors, the tape of M_1 is split into *tracks*: each track corresponds to a *field* of a cell symbol of M_1 viewed as a data record. Each stage of computation will be repeated three times. The results will be stored in separate tracks, and a final cell-by-cell majority vote will read out the result of the work period from them.

All this organization is controlled by a few key fields, for example a field called *Addr* showing the position of each cell in the colony, and a field *Sweep* showing the last sweep of the computation (along with its direction) that has been performed already. The most technical part is to protect this control information from bursts.

For example, a burst can reverse the head in the middle of a sweep. We want to discover such structural disruptions locally, so the head will not be allowed to go far from the place where its sweep was turned, without looking back. Therefore the head will make frequent zigzags even during a single sweep. This will trigger the healing procedure if for example a premature turn-back is detected.

Note that the healing procedure also must resist interruption (or false start) by a burst.

Remark 2.1 This description uses some words in an informal way. Some of these will get precise definition later. For example, the word *colony* will have at least two formal definitions. From the point of view of the program (transition function), it is just a sequence of addresses from 0 to $Q - 1$. From the point

of view of the analysis, it is a sequence of actual adjacent tape cells with the address field having these values. \lrcorner

2.2 Hierarchical construction

In order to build a machine resisting faults occurring independently of each other with some small probability, we take the approach suggested in [5], and implemented in [3] and [4] for the case of one-dimensional cellular automata, with some ideas from the tiling application of [2]. We will build a *hierarchy of simulations*: machine M_1 simulates machine M_2 which simulates machine M_3 , and so on. To simplify the outline here, assume now that all these machines have the same program, and all simulations have the same block size. (Later, each level k will have its own cell size B_k and block size Q_k , with $B_{k+1} = B_k Q_k$.)

One cell of machine M_3 is simulated by one colony of machine M_2 . Correspondingly, one cell of M_2 is simulated by one colony of machine M_1 . So one cell of M_3 is simulated by Q^2 cells of M_1 . Further, one step of machine M_3 is simulated by one work period of M_2 of, say, $O(Q^2)$ steps. One step of M_2 is simulated by one work period of M_1 , so one step of M_3 is simulated by $O(Q^4)$ steps of M_1 .

Per construction, machine M_1 can withstand bursts of faults with size $\leq \beta$ for some constant parameter β , separated by some $O(Q^2)$ fault-free steps. Machines M_2, M_3, \dots have the same program, so it would be natural to expect that machine M_1 can withstand also some *additional*, larger bursts of size $\leq \beta Q$ if those are separated by at least $O(Q^4)$ steps.

But a new obstacle arises. Damage caused by a big burst spans several colonies. The repair mechanism of machine M_1 outlined in Section 2.1 above is too local to recover from such extensive damage. This cannot be allowed, since then the whole hierarchy would stop working. So we add a new mechanism to M_1 that, more modestly, will just try to restore a large enough portion of the tape (of the extent of several colonies), so it can go on with the simulation of M_2 . Note that this mechanism does not aim to restore the lost original information: that job is left to higher levels. It just tries to restore enough structure to enable the (simulated) work of M_2 .

All machines above M_1 in the hierarchy are “virtual”: the only hardware in the construction is machine M_1 . Moreover, they will not be ordinary Turing machines, but *generalized* ones, with some new features that are not needed on the lowest level but seem necessary in a simulated Turing machine: for example they allow a positive distance between neighboring tape cells.

A tricky issue is “forced self-simulation”: while we are constructing machine

M_1 we want to give it the feature that it will simulate a machine M_2 that works just like M_1 . The “forced” feature means that this simulation should work without any written program (since that could be corrupted by faults).

This will be achieved by a construction similar to the proof of the Kleene’s fixed-point theorem (also called recursion theorem). We first fix a (simple) programming language to express the transition function of a Turing machine. We write an interpreter for it in this same language (just as compilers for the C language are sometimes written in C). The program of the transition function of M_2 (essentially the same as that of M_1) in this language, is a string that will be “hard-wired” into the transition function of M_1 , so that the latter, at the start of each work period, can write it on a working track of the current colony. Then the work period will interpret it, applying it to the data found there, resulting in the simulation of M_2 .

This gives rise to an infinite sequence of simulations, to withstand larger and larger but sparser and sparser bursts of faults.

Since the M_1 uses the universal interpreter, which in turns simulates the same program, it is natural to ask how machine M_1 simulates a given Turing machine G that does the actual useful computation? For this task, we set aside a separate track on each machine M_i , on which some arbitrary other Turing machine can be simulated. The higher the level of the machine M_k that performs this “side-simulation”, the higher the reliability. Thus, only the simulations $M_k \rightarrow M_{k+1}$ are forced, without program (that is with a “hard-wired” program): the side simulations can rely on written programs, since the robust structure in the hierarchy M_1, M_2, \dots will support them.

2.3 From combinatorial to probabilistic noise

The construction outlined above was related to burst increasing in size and decreasing in frequency. In essence, that noise model is combinatorial. To deal with probabilistic noise combinatorially, we stratify the set of faulty times *Noise* as follows. For a series of parameters β_k, V_k , we first remove “isolated bursts” of type (β_1, V_1) of elements of this set. (See below the notion of “isolated bursts” of type (β, V) .) Then, we remove isolated bursts of type (β_2, V_2) from the remaining set, and so on. It will be shown that with the appropriate choice of parameters, with probability 1, eventually nothing is left over from the set *Noise*.

A composition of two reliable simulations is even more reliable. We will see that a sufficiently large hierarchy of such simulations resists probabilistic noise.

2.4 Difficulties

We list here some of the main problems that the paper deals with, and some general ways they will be solved or avoided. Some more specific problems will be pointed out later, along with their solution.

Non-aligned colonies A large burst of M_1 can modify the order of entire colonies or create new ones with gaps between them.

To overcome this problem conceptually, we introduce the notion of a *generalized Turing machine* allowing for non-adjacent cells. Each such machine has a parameter B called the *cell body size*. The cell body size of a Turing machine in Section 1.2 would still remain 1.

No structure What to do when the head is in a middle of an empty area where no structure exists? To ensure reliable passage across such areas, we will try to keep everything filled with cells, even if these are not part of the main computation.

Clean areas Noise can create areas over which the predictability of the simulated machine is limited. In these areas the structure of the underlying simulation (invisible on this level) may be destroyed. These areas should not simply be considered blank, since blankness implies predictable behavior. We call these areas of “disordered”, and call the complement *clean*. Transition properties will allow making conclusions from cleanness, not from disorder. Examples below will show that when the head enters disorder, it can be “sucked in”.

Extending cleanness The definition of the generalized Turing machine will stipulate some “magical” properties helping to restore cleanness:

- A) extension of clean intervals as the head passes in and out of them;
- B) the appearance of a clean “hole” around the head whenever it passes a certain amount of time in a small interval;
- C) the cleaning of an interval if it is passed noiselessly a certain number of times.

(The need to implement these same properties in simulation is one of the main burdens of the construction.) While an area is cleaned, it will also be re-populated with cells. Their content is not important, what matters is the restoration of predictability.

Rebuilding If local repair fails, a special rule will be invoked that reorganizes a larger part of the tape (of the size of a few colonies instead of only a few cells). This is the mechanism implementing the above “magical” properties on the next level.

The following examples show the difficulties that necessitate some of the complexities of the construction and proof. Some of them may not be completely understandable without the details of the following program. You can skip these examples safely, and return to them later when you are wondering about the motivation for some feature.

Example 2.2 (Need for zigging) When the head works in a colony of machine M_1 performing a simulation of a cell of machine M_2 , it works in sweeps across the colony. But a burst of faults could reverse the head in the middle of a sweep, leaving it uncompleted. This way a local burst could create non-local inconsistency. ┘

To handle the problem of Example 2.2, the head will proceed in zigzags: every step advancing the head in the simulation is followed by Z steps of going backward and forward again, just checking consistency (and starting a healing process if necessary). The parameter Z will be chosen appropriately.

Example 2.3 (Need for feathering) Some big noise can create a number of intervals I_1, I_2, \dots, I_n consisting of colonies of machine M_1 , each interval with its own simulated head, where the neighboring intervals are in no relation to each other. When the head is about to return from the end of I_k (never even to zig beyond it, see the discussion after Example 2.2), a burst can carry it over to I_{k+1} where the situation may be symmetric: it will continue the simulation that I_{k+1} is performing. (The rightmost colony of I_k and the leftmost colony of I_{k+1} need not be complete: what matters is only that the simulation in I_k would not bring the head beyond its right end, and the simulation in I_{k+1} would not bring the head beyond its left end.)

The head can be similarly captured to I_{k+2} , then much later back from I_{k+1} to I_k , and so on. This way the restoration of structure in M_2 may be delayed too long. ┘

The device by which we will mitigate the effect of this kind of capturing is another property of the movement of the head which will call *feathering*: if the head turns back from a tape cell then next time it must go beyond. This requires a number of adjustments to the program (see later).

Example 2.4 (Two slides over disorder) This example shows the possibility for the head to slide twice over disorder without cleaning it.

Consider two levels of simulation as outlined in Section 2.2: machine M_1 simulates M_2 which simulates M_3 . The tape of M_1 is subdivided into colonies of size Q_1 . A burst on level 1 has size $O(1)$, while a burst on level 2 has size $O(Q_1)$.

Suppose that M_1 is performing a simulation in colony C_0 . An earlier big burst may have created a large interval D of disorder on the right of C_0 , even reaching into C_0 . For the moment, let C_0 be called a *victim* colony. Assume that the left edge of D represents the last stage of a transfer operation to the right neighbor colony $C_0 + Q_1$. When the head, while performing its work in C_0 , moves close to its right end, a small burst may carry it over into D . There it will be “captured”, and continue the (unintended) right transfer operation. This can carry the head, over several successful colony simulations in D , to some victim colony C_1 on the right from which it will be captured to the right similarly. This can continue over new and new victim colonies C_i (with enough space between them to allow for new bursts to occur), all the way inside the disorder D . So the M_2 cells in D will fail to simulate M_3 .

After a while the head may return to the left in D (performing the simulations in its colonies). When it gets at the right end of a victim colony C_i , a burst might move it back there. There is a case when C_i now can just continue its simulation and then send the head further left: when before the head was captured on its right, it was in the last stage of simulating a left turn of the head of machine M_2 .

In summary, a big burst can create a disordered area D which can capture the head and on which the head can slide forward and back without recreating any level of organization beyond the second one. \lrcorner

Example 2.5 (Many slides over disorder) Let us describe a certain “organization” of a disordered area in which an unbounded number of passes may be required to restore order. For some $n < 0$, let the cells of M_1 at positions x_{-Q_1}, \dots, x_n , where $x_{i+1} = x_i + B_1$, represent part of a healthy colony $C(x_{-Q_1})$ starting at x_{-Q_1} , where x_n is the rightmost cell of $C(x_{-Q_1})$ to which the head would come in the last sweep before the simulation will move to the *left* neighbor colony $C(x_{-2Q_1})$. Let them be followed by cells $x_{n+1}, \dots, x_{Q_1-1}, \dots$ which represent the last sweep of a transfer operation to the *right* neighbor colony $C(x_0)$. If the head is in cell x_n , a burst can transfer it to x_{n+1} . The cell state of M_2 simulated by $C(x_{-Q_1})$ need to be in *no relation* to the cell state of M_2 simulated by $C(x_0)$. This was a capture of the head by a burst of M_1 across the point 0, to the right.

We can repeat the capture scenario, say around points iQ_1Q_2 for $i = 1, 2, \dots$, and this way cells of M_3 simulated by M_2 (simulated by M_1) can be defined arbitrarily, with no consistency needed between any two neighbors. (We did not write iQ_1 just in case bursts are not allowed in neighboring colonies.) In particular, we can define them to implement a *leftward* capture scenario via level 3 bursts at points $iQ_1Q_2Q_3Q_4$, allowing to simulate arbitrary cells of M_5 with

no consistency requirement between neighbors. So M_5 could again implement a rightward capture scenario, and so on. In summary, a malicious arrangement of disorder and noise allows k passes after which the level of organization is still limited to level $2k + 1$. \lrcorner

Our construction will ensure that, on the other hand, $O(k)$ passes (free of k -level will restore organization to level k . This property of the construction will be incorporated into our definition of a generalized Turing machines as the “magical” property (C) above.

3 Notation

Most notational conventions given here are common; some other ones will also be useful.

Natural numbers and integers By \mathbb{Z} we denote the set of integers.

$$\begin{aligned}\mathbb{Z}_{>0} &= \{x : x \in \mathbb{Z}, x > 0\}, \\ \mathbb{Z}_{\geq 0} &= \mathbb{N} = \{x : x \in \mathbb{Z}, x \geq 0\}.\end{aligned}$$

Intervals We use the standard notation for intervals:

$$\begin{aligned}[a, b] &= \{x : a \leq x \leq b\}, & [a, b) &= \{x : a \leq x < b\}, \\ (a, b] &= \{x : a < x \leq b\}, & (a, b) &= \{x : a < x < b\}.\end{aligned}$$

We will also write $[a, b)$ in place of $[a, b) \cap \mathbb{Z}$, whenever this leads to no confusion. Instead of $[x + a, x + b)$, sometimes we will write

$$x + [a, b).$$

Ordered pairs Ordered pairs are also denoted by (a, b) , but it will be clear from the context whether we are referring to an ordered pair or open interval.

Comparing the order of a number and an interval For a given number x and interval I , we write

$$x \geq I$$

if for every $y \in I$, $x \geq y$.

Distance The distance between two real numbers x and y is defined in a usual way:

$$d(x, y) = |x - y|.$$

The distance of a point x from interval I is

$$d(x, I) = \min_{y \in I} d(x, y).$$

Ball, neighborhood, ring, stripe A ball of radius $r > 0$, centered at x is

$$B(x, r) = \{y : d(x, y) \leq r\}.$$

An r -neighborhood of interval I is

$$\{x : d(x, I) \leq r\}.$$

An r -ring around interval I is

$$\{x : d(x, I) \leq r \text{ and } x \notin I\}.$$

An r -stripe to the right of interval I is

$$\{x : d(x, I) \leq r \text{ and } x \notin I \text{ and } x > I\}.$$

Logarithms Unless specified differently, the base of logarithms throughout this work is 2.

4 Specifying a Turing machine

Let us introduce the tools allowing to describe the reliable Turing machine.

4.1 Universal Turing machine

We will describe our construction in terms of universal Turing machines, operating on binary strings as inputs and outputs. We define universal Turing machines in a way that allows for rather general “programs”.

Definition 4.1 (Standard pairing) For a (possibly empty) binary string $x = x(1) \cdots x(n)$ let us introduce the map

$$\langle x \rangle = 0^{|x|}1x,$$

Now we encode pairs, triples, and so on, of binary strings as follows:

$$\begin{aligned} \langle s, t \rangle &= \langle s \rangle t, \\ \langle s, t, u \rangle &= \langle \langle s, t \rangle, u \rangle, \end{aligned}$$

and so on.

From now on, we will assume that our alphabet Σ is of the form $\Sigma = \{0, 1\}^s$, that is the tape symbols are viewed as binary strings of a certain length. Also, if we write $\langle i, u \rangle$ where i is some number, it is understood that the number i is represented in a standard way by a binary string. \lrcorner

Definition 4.2 (Computation result, universal machine) Assume that a Turing machine M starting on binary x , at some time t writes the first time the blank symbol at position 0 (this is interpreted as halting). Then we look at the longest (possibly empty) binary string to be found starting at position 1 on the tape, and call it the *computation result* $M(x)$. We will write

$$M(x, y) = M(\langle x, y \rangle), \quad M(x, y, z) = M(\langle x, y, z \rangle),$$

and so on.

A Turing machine U is called *universal* among Turing machines with binary inputs and outputs, if for every Turing machine M , there is a binary string p_M such that for all x we have $U(p_M, x) = M(x)$. (This equality also means that the computation denoted on the left-hand side halts if and only if the computation on the right-hand side does.) \lrcorner

Let us introduce a special kind of universal Turing machines, to be used in expressing the transition functions of other Turing machines. These are just the Turing machines for which the so-called s_{mn} theorem of recursion theory holds with $s(x, y) = \langle x, y \rangle$.

Definition 4.3 (Flexible universal Turing machine) A universal Turing machine will be called *flexible* if whenever p has the form $p = \langle p', p'' \rangle$ then

$$U(p, x) = U(p', \langle p'', x \rangle).$$

Even if x has the form $x = \langle x', x'' \rangle$, this definition chooses $U(p', \langle p'', x \rangle)$ over $U(\langle p, x' \rangle, x'')$, that is starts with parsing the first argument (this process converges, since x is shorter than $\langle x, y \rangle$). \lrcorner

It is easy to see that there are flexible universal Turing machines. On input $\langle p, x \rangle$, a flexible machine first checks whether its “program” p has the form $p = \langle p', p'' \rangle$. If yes, then it applies p' to the pair $\langle p'', x \rangle$. (Otherwise it just applies p to x .)

Definition 4.4 (Transition program) Consider an arbitrary Turing machine M with alphabet Σ , and transition function τ . A binary string π will be called a

transition program of M if whenever $\tau(\mathbf{a}) = (\mathbf{a}', j)$ we have

$$U(\pi, \mathbf{a}) = \langle \mathbf{a}, j \rangle.$$

We will also require that the computation induced by the program makes $O(|\pi| + |a|)$ left-right turns, over a length tape $O(|\pi| + |a|)$. \lrcorner

In our simulations of a Turing machine M_2 on some other machine M_1 , the transition program of M_2 will just provide a way to compute the (local) transition function of M_2 by the universal machine; it does not organize the rest of the simulation.

Remark 4.5 In the construction of universal Turing machines provided by the textbooks (though not in the original one given by Turing), the program is generally a string encoding a table for the transition function τ of the simulated machine M . Other types of program are imaginable: some simple transition functions can have much simpler programs. However, our fixed machine is good enough (similarly to the optimal machine for Kolmogorov complexity). If some machine U' simulates M via a very simple program q , then

$$M(x) = U'(q, x) = U(p_{U'}, \langle q, x \rangle) = U(\langle p_{U'}, q \rangle, x),$$

so U simulates this computation via the program $\langle p_{U'}, q \rangle$. \lrcorner

4.2 Rule language

In what follows we will describe the generalized Turing machines M_k for all k . They are all similar, differing only in the parameter k ; the most important activity of M_k is to simulate M_{k+1} . The description will be uniform, except for the parameter k . We will denote therefore M_k simply by M , and M_{k+1} by M^* . Similarly we will denote the block size Q_k of the block code of the simulation simply by Q .

Instead of writing a huge table describing the transition function $\tau_k = \tau$, we present the transition function as a set of *rules*. It will be then possible to write one *interpreter* program that carries out these rules; that program can be written for some fixed flexible universal machine Univ.

Each rule consists of some (nested) conditional statements, similar to the ones seen in an ordinary program: “**if** condition **then** instruction **else** instruction”, where the condition is testing values of some fields of the observed cell pair, and the instruction can either be elementary, or itself a conditional statement. The elementary instructions are an *assignment* of a value to a field of a cell symbol,

or a command to move the head. Rules can call other rules, but these calls will never form a cycle. Calling other rules is just a shorthand for nested conditions.

Even though rules are written like procedures of a program, they describe a single transition. When several consecutive statements are given, then they change different fields of the state or cell symbol, so they can be executed simultaneously.

Assignment of value x to a field y of the state or cell symbol will be denoted by $y \leftarrow x$. We will also use some conventions introduced by the C language: namely, $x \leftarrow x + 1$ and $x \leftarrow x - 1$ are abbreviated to $x++$ and $x--$ respectively.

Rules can also have parameters, like $\text{Swing}(a, b, u, v)$. Since each rule is called only a constant number of times in the whole program, the parametrized rule can be simply seen as a shorthand.

Mostly we will describe the rules using plain English, but it should always be clear that they are translatable into such rules.

4.3 Fields

For the machine M we are constructing, each tape symbol will be a tuple $q = (q_1, q_2, \dots, q_k)$, where the individual elements of the tuple will be called *fields*, and will have symbolic names. For example, we will have fields *Addr* and *Drift*, and may write q_1 as $q.\text{Addr}$ or just *Addr*, q_2 as $q.\text{Drift}$ or *Drift*, and so on.

The array of values of the same field of the cells will be called a *track*. Thus, we will talk about the *Hold* track of the tape, carrying the *Hold* fields of the cells. Each field of a cell has a possible value \emptyset whose approximate meaning is “undefined”.

In what follows we describe some of the most important fields we will use; others will be introduced later.

A properly formatted configuration of M splits the tape into blocks of Q consecutive cells called *colonies*. One colony of the tape of the simulating machine represents one cell of the simulated machine. The two colonies that correspond to the pair of cells that the simulated machine is scanning is called the *base colony-pair* (a precise definition will refer to the actual history of the work of M), and its members will be called the *left* and *right base colony*. Sometimes the left base colony will just be called the *base colony*. Most of the computation proceeds over the base colony-pair. The direction of the simulated head movement, once figured out by the computation, is called the *drift*. Two neighbor colonies may not be adjacent, in which case the cells will form a *bridge* between them.

The present behavior of the computation will be characterized by a field called the *mode*:

$$Mode \in \{\text{Normal, Healing, Rebuilding}\}.$$

If its value in the current cell is Normal we will say that the computation is in *normal* mode. In this case, the machine is engaged in the regular business of simulation. The *healing* mode tries to correct some local fault due to a couple of neighboring bursts, while the *rebuilding* mode attempts to restore the colony structure on the scale of a couple of colonies. The

Info

track of a colony of M contains the string that encodes the content of the simulated cell of M^* . The

Prog

track stores the program of M^* , in an appropriate form to be interpreted by the simulation. (As mentioned in Section 2.2, this field will be subject to some “hard-wiring” manipulation.) The field

Addr

of the cell shows the position of the cell in its colony: it takes values in $[0, Q)$. The direction in $\{-1, 1\}$ in which the simulated head moves will be recorded on the track

Drift.

The

Sweep

field counts the sweeps that the head makes during the work period. The number of the last sweep of the work period will depend on the drift d , and will be denoted by

$$\text{Last}(d). \tag{4.1}$$

In calculating parameters, we will make use of

$$V = \max(\text{Last}(-1), \text{Last}(1)). \tag{4.2}$$

Cells will be designated as belonging to a number of possible *kinds*, signaled by the field

Kind

with values Vac, Stem, Member₀, Member₁, Bridge. Here is a description of the role of these cell kinds.

Vac is not really the kind of any cell, it is just used in the transition function to show that there is no cell or to describe the action of killing a cell. Stem is the kind of newly created cells that do not participate in any known colony structure. We will also try to keep all areas between colonies filled with (not necessarily adjacent) stem cells. For example the computation may find that a colony does not properly encode a tape cell by the required error-correcting code. Then we want to “kill” the whole colony. This will happen by turning the kind of each of its cells to Stem.

Cells of the base colony-pair are of type Member₀ and Member₁ respectively. If the base colony pair is not adjacent then there will be a *bridge* of $< Q$ adjacent cells of type Bridge between them.

Definition 4.6 We will say that the bridge *extends* a colony if its cells are adjacent to it. J

Under normal conditions, if the *Drift* track on the bridge has value -1 then it is extending the left base colony, otherwise is extending the right base colony. Cells of the bridge will continue the addressing (modulo Q) of the colony it extends.

During healing, some special fields of the state and cell are used, treated as subfields of the field

Heal. (4.3)

In particular, there will be a *Heal.Sweep* field. During rebuilding, we will work with subfields of the field

Rebuild,

and a cell will be called *marked for rebuilding* if *Rebuild.Sweep* $\neq 0$.

5 Exploiting structure in the noise

5.1 Sparsity

Let us introduce a technique connecting the combinatorial and probabilistic noise models.

Definition 5.1 (Centered rectangles, isolation) Let $\mathbf{r} = (r_1, r_2)$, $r_1, r_2 \geq 0$, be a two-dimensional nonnegative vector. A *rectangle* of radius \mathbf{r} centered at \mathbf{x} is

$$B(\mathbf{x}, \mathbf{r}) = \{\mathbf{y} : |y_i - x_i| \leq r_i, i = 1, 2\}. \quad (5.1)$$

Let $E \subseteq \mathbb{Z}^2$ be a two-dimensional set. A point \mathbf{x} of E is $(\mathbf{r}, \mathbf{r}^*)$ -isolated if

$$E \cap B(\mathbf{x}, \mathbf{r}^*) \subseteq B(\mathbf{x}, \mathbf{r}).$$

Set E is $(\mathbf{r}, \mathbf{r}^*)$ -sparse if $D(E, \mathbf{r}, \mathbf{r}^*) = \emptyset$, that is it consists of $(\mathbf{r}, \mathbf{r}^*)$ -isolated points. Let

$$D(E, \mathbf{r}, \mathbf{r}^*) = \{\mathbf{x} \in E : \mathbf{x} \text{ is not } (\mathbf{r}, \mathbf{r}^*)\text{-isolated from } E\}. \quad (5.2)$$

┘

Definition 5.2 (Sparsity) Let

$$\beta \geq 9 \quad (5.3)$$

be a parameter, and let

$$0 < B_1 < B_2 < \cdots, \quad T_1 < T_2 < \cdots, \quad 1 \leq \gamma_1 \leq \gamma_2 \leq \cdots$$

be sequences of positive integers to be fixed later. For a two-dimensional set E , let $E^{(1)} = E$. For $k > 1$ we define recursively:

$$E^{(k+1)} = D(E^{(k)}, \beta(B_k, T_k), \gamma_k(B_{k+1}, T_{k+1})). \quad (5.4)$$

Set $E^{(k)}$ is called the k -th residue of E . It is k -sparse if $E^{(k+1)} = \emptyset$. It is simply sparse if $\bigcap_k E^{(k)} = \emptyset$.

When $E = E^{(k)}$ and k is known then we will denote $E^{(k+1)}$ simply by E^* . ┘

The following lemma connects the above defined sparsity notions to the requirement of small fault probability. It is formulated somewhat redundantly, for easier application.

Lemma 5.3 (Sparsity) *Let $Q_k = B_{k+1}/B_k$, $U_k = T_{k+1}/T_k$, and*

$$\lim_{k \rightarrow \infty} \frac{\log(\gamma_k U_k Q_k)}{1.5^k} = 0. \quad (5.5)$$

For sufficiently small ε , for every $k \geq 1$ the following holds. Let $E \subseteq \mathbb{Z} \times \mathbb{Z}_{\geq 0}$ be a random set with the property that each pair (p, t) belongs to E independently from the other ones with probability $\leq \varepsilon$.

Then for each point \mathbf{x} and each k ,

$$\mathbb{P}\{B(\mathbf{x}, (B_k, T_k)) \cap E^{(k)} \neq \emptyset\} < 2\varepsilon \cdot 2^{-1.5^k}.$$

As a consequence, the set E is sparse with probability 1.

Proof. Let $k = 1$. Rectangle $B(\mathbf{x}, (B_1, T_1))$ is a single point, hence the probability of our event is $< \varepsilon$. Let us prove the inequality by induction, for $k + 1$.

Note that our event depends at most on the rectangle $B(\mathbf{x}, 3(B_k, T_k))$. Let

$$p_k = 2\varepsilon \cdot 2^{-1.5^k}.$$

Suppose $\mathbf{y} \in E^{(k)} \cap B(\mathbf{x}, \gamma_k(B_{k+1}, T_{k+1}))$. Then, according to the definition of $E^{(k)}$, there is a point

$$\mathbf{z} \in B(\mathbf{y}, \gamma_k(B_{k+1}, T_{k+1})) \cap E^{(k)} \setminus B(\mathbf{y}, \beta(B_k, T_k)). \quad (5.6)$$

Consider a standard partition of the (two-dimensional) space-time into rectangles $K_p = \mathbf{c}_p + [-B_k, B_k] \times [-T_k, T_k]$ with centers $\mathbf{c}_1, \mathbf{c}_2, \dots$. The rectangles K_i, K_j containing \mathbf{y} and \mathbf{z} respectively intersect $B(\mathbf{x}, 2\gamma_k(B_{k+1}, T_{k+1}))$. The triple-size rectangles $K'_i = \mathbf{c}_i + [-3B_k, 3B_k] \times [-3T_k, 3T_k]$ and K'_j are disjoint, since (5.3) and (5.6) imply $|\gamma_1 - \gamma_2| > \beta B_k$ and $|\gamma_1 - \gamma_2| > \beta T_k$.

The set $E^{(k)}$ must intersect two rectangles K_i, K_j of size $2(B_k, T_k)$ separated by at least $4(B_k, T_k)$, of the big rectangle $B(\mathbf{x}, 2\gamma_k(B_{k+1}, T_{k+1}))$.

By the inductive hypothesis, the event \mathcal{F}_i that K_i intersects E_k has probability bound p_k . It is independent of the event \mathcal{F}_j , since these events depend only on the triple size disjoint rectangles K'_i and K'_j .

The probability that both of these events hold is at most p_k^2 . The number of possible rectangles K_p intersecting $B(\mathbf{x}, 2\gamma_k(B_{k+1}, T_{k+1}))$ is at most $C_k := ((2\gamma_k^2 U_k Q_k) + 2)^2$, so the number of possible pairs of rectangles is at most $C_k^2/2$, bounding the probability of our event by

$$\begin{aligned} C_k^2 p_k^2 / 2 &= 2C_k^2 \varepsilon^2 2^{-1.5^{k+1}} \cdot 2^{-0.5 \cdot 1.5^k} \\ &= 2\varepsilon 2^{-1.5^{k+1}} \cdot \varepsilon C_k^2 2^{-0.5 \cdot 1.5^k}. \end{aligned}$$

Since $\lim_k \frac{\log(\gamma_k U_k Q_k)}{1.5^k} = 0$, the last factor is ≤ 1 for sufficiently small ε . \square

In our construction we will choose

$$\gamma_k = \Omega(k^2), \quad Q_k = \Omega(\gamma_k^2), \quad U_k = \Omega(Q_k^2),$$

while still satisfying (5.5).

5.2 Error-correcting code

Let us add error-correcting features to the block codes introduced in Definition 1.2.

Definition 5.4 (Error-correcting code) A block code is (β, t) -burst-error-correcting, if for all $x \in \Sigma_2$, $y \in \Sigma_1^Q$ we have $\psi^*(y) = x$ whenever y differs from $\psi_*(x)$ in at most t intervals of size $\leq \beta$. For such a code, we will say that a word $y \in \Sigma_1^Q$ is r -compliant if it differs from a codeword of the code by at most r intervals of size $\leq \beta$. \lrcorner

Example 5.5 (Repetition code) Suppose that $Q \geq 3\beta$ is divisible by 3, $\Sigma_2 = \Sigma_1^{Q/3}$, $\psi_*(x) = xxx$. If $y = y(1) \dots y(Q)$, then $x = \psi^*(y)$ is defined by $x(i) = \text{maj}(y(i), y(i + Q/3), y(i + 2Q/3))$. For all $\beta \leq Q/3$, this is a $(\beta, 1)$ -burst-error-correcting code. If we repeat 5 times instead of 3, we get a $(\beta, 2)$ -burst-error-correcting code. Let us note that there are much more efficient such codes than just repetition. \lrcorner

Consider a Turing machine (Σ, τ) (actually a generalized one to be defined later) simulating some Turing machine (Σ^*, τ^*) . We will assume that the alphabet Σ^* is a subset of the set of binary strings $\{0, 1\}^l$ for some $l < Q$ (we can always ignore some tape symbols, if we want).

Definition 5.6 (Interior) Let

$$PadLen$$

be a parameter to be defined later (in (7.5)). Consider an interval I of neighbor cells. A cell belongs to the *interior* of I if there are at least $PadLen$ neighbors between it and the complement of I . In particular, we will talk about the interior of a colony. \lrcorner

We will store the coded information in the interior of the colony, since it is more exposed to errors near the boundaries. So let (v_*, v^*) be a $(\beta, 2)$ -burst-error-correcting block code with

$$v_* : \{0, 1\}^l \cup \{\emptyset\} \rightarrow \{0, 1\}^{(Q-2 \cdot PadLen)B}.$$

We could use, for example, the repetition code of Example 5.5. Other codes are also appropriate, but we require that they have some fixed programs p_{encode} , p_{decode} on the universal machine Univ, in the following sense:

$$v_*(x) = \text{Univ}(p_{\text{encode}}, x), \quad v^*(y) = \text{Univ}(p_{\text{decode}}, y).$$

Also, these programs must work in quadratic time and linear space on a one-tape Turing machine (as the repetition code certainly does).

Let us now define the block code (ψ_*, ψ^*) used below in the definition of the configuration code (φ_*, φ^*) outlined in Section 6.3:

$$\psi_*(a) = 0^{\text{PadLen}} v_*(a) 0^{\text{PadLen}}. \quad (5.7)$$

The decoded value $\psi^*(x)$ is obtained by first removing PadLen symbols from both ends of x to get x' , and then computing $v^*(x')$. It will be easy to compute the configuration code from ψ_* , once we know what fields there need initialization.

6 The model

Recall the definition of sparsity in Section 5.1: there will be a sequence $0 < B_1 < B_2 < \dots$ of “scales” in space and a sequence $T_1 < T_2 < \dots$ of scales in time, and a constant β . We will define a sequence of simulations $M_1 \rightarrow M_2 \rightarrow \dots$ where each M_k is a machine simulating one on a higher level. For simplicity, we will use the notation $M = M_k$, $M^* = M_{k+1}$, and similarly for the other parameters, for example B, T, Q, U . As indicated in Section 2.2, these machines will generalize ordinary Turing machines, with a number of new features.

Notation 6.1 For an interval $I = [a, b)$ and some $c \geq 0$ let

$$\text{Int}(I, c) = [a + c, b - c).$$

Similarly for intervals $(a, b]$, $[a, b]$, (a, b) . ┘

6.1 Generalized Turing machine

Standard Turing machines do not have operations like “creation” or “killing” of cells, nor do they allow for cells to be non-adjacent. We introduce here a *generalized Turing machine*. It depends on an integer $B \geq 1$ that denotes the cell body size, and an upper bound T on the transition time. These parameters provide the illusion that the different Turing machines in the hierarchy of simulations all operate on the same linear space and the same time interval. Even if the notions

of cells and alphabet are different for each machine of the hierarchy, at least the notion of a *location on the tape* is the same. There will also be a *pass number* parameter π , whose meaning will be explained in the definition of trajectories.

Definition 6.2 (Generalized Turing machine) A *generalized Turing machine* M is defined by a tuple

$$(\Sigma, \tau, Adj, B, T, \pi), \quad (6.1)$$

where Σ is a finite set called the *alphabet*,

$$\tau : \Sigma^2 \rightarrow \Sigma^2 \times \{-1, 1\}$$

is the *transition function*, Adj is a function on the alphabet (can be called a “field” if a symbol is viewed as a piece of data, a record with several fields). $Adj : \Sigma \rightarrow \{0, 1\}$ will show whether the last move to the present cell was from an adjacent cell (it could have been for example from a non-adjacent neighbor). The integer $B \geq 1$ is called the *cell body size*, and the real number T is a bound on the transition time. The integer π will be the number of passes needed to clean an area (see below). Among the elements of the tape alphabet Σ , we distinguish the elements 0, 1, Bad, Vac. The role of the symbols Bad and Vac will be clarified below. (They can be viewed to correspond to the “blank” symbol of a regular Turing machine.) \lrcorner

The definition of a configuration below introduces a new concept, the *current cell-pair*, a pair of positions $\hat{\mathbf{h}} = (\hat{h}_0, \hat{h}_1)$. (We will call \hat{h}_0 the *current cell*.) The current cell’s position may not exactly be the position of the head: here is some explanation. A generalized Turing machine M^* may be simulated by some lower-level (possibly generalized) Turing machine M . Not all details of the tape content of M are visible on the level of M^* : what is visible is the content of the cells of M^* decoded from those of M . The currently observed cell-pair in the simulated machine M^* is $\hat{\mathbf{h}}$. On the other hand the true head position h of M may perform a lot of oscillatory movement within and around the current cell pair not directly relevant to the work of M^* .

Definition 6.3 (Configuration) Consider a generalized Turing machine (6.1). A *configuration* is a tuple

$$(A, h, \hat{\mathbf{h}}),$$

where $A : \mathbb{Z} \rightarrow \Sigma$, $h, \hat{h}_j \in \mathbb{Z}$. The array A is the tape configuration. It is Vac in all but finitely many positions.

As in (1.1) before, h is the head position, but it may differ from the current cell \hat{h}_0 , so now if $\xi = (A, h, \hat{\mathbf{h}})$ then we write

$$\xi.\text{tape} = A, \quad \xi.\text{pos} = h, \quad \xi.\text{cur-cell}_j = \hat{h}_j. \quad (6.2)$$

A point p is *clean* if $A(p) \neq \text{Bad}$. A set of points is *clean* if it consists of clean points. Whenever the interval $h + [4B, 4B)$ is clean the current cell-pair must be within it.

We say that there is a *cell* at a position $p \in \mathbb{Z}$ if the interval $p + [0, B)$ is clean and $A(p) \neq \text{Vac}$. In this case, we call the interval $p + [0, B)$ the *body* of this cell. Cells must be at distance $\geq B$ from each other, that is their bodies must not intersect. If they are at a distance $< 2B$ from each other then they are called *neighbors*. They are called *adjacent* if the distance is exactly B . A configuration will always have the following property:

C1) Any endpoint of a maximal clean interval is the end of a cell body in that direction.

Let

$$\text{Configs}_M$$

denote the set of all possible configurations of a Turing machine M . \lrcorner

All the above definitions can clearly be localized to define a configuration *over a space interval* I , where it is always understood that $h \in I$, that is I contains the head.

Definition 6.4 (Local configuration, replacement) A *local configuration* on a (finite or infinite) interval I is given by values assigned to the cells of I , along with the following information: whether the head is to the left of, to the right of or inside I , and if it is inside, on which cell.

If I' is a subinterval of I , then a local configuration ξ on I clearly gives rise to a local configuration $\xi(I')$ on I' as well, called its *subconfiguration*. We can similarly talk about local tape configurations.

Let ξ be a tape configuration and $\zeta(I)$ a local tape configuration. Then the tape configuration $\xi|\zeta(I)$ is obtained by replacing ξ with ζ over the interval I . \lrcorner

It is natural to name a sequence of configurations that is conceivable as a computation (faulty or not) of a Turing machine as “history”. The histories that obey the transition function then could be called “trajectories”. In the following definition we will stretch this notion to encompass also some limited violations

of the transition function. In connection with any underlying Turing machine with a given starting configuration, we will denote by

$$\text{Noise} \subseteq \mathbb{Z} \times \mathbb{Z}_{\geq 0} \quad (6.3)$$

the set of space-time points (p, t) , such that a fault occurs at time t when the head is at position p .

Definition 6.5 (History) For a generalized Turing machine (6.1), consider a sequence $\eta = (\eta(0), \eta(1), \dots)$ of configurations with $\eta(t) = (A(\cdot, t), h(t), \hat{h}(t))$, along with a noise set Noise . The *switching times* of this sequence are the times t when the position or content of the current cell-pair changes. The interval between two consecutive switching times is the *dwell period*. The pair

$$(\eta, \text{Noise})$$

will be called a *history* of machine M if the following conditions hold.

- $|h(t) - h(t')| \leq |t' - t|$.
- In two consecutive configurations, the content $A(p, t)$ of the positions p not in $h(t) + [-2B, 2B]$, remains the same: for example $A(n, t + 1) = A(n, t)$ for all $n \notin h(t) + [-2B, 2B]$.
- At each noise-free switching time the head is on the new current cell: $\hat{h}_0(t) = h(t)$. In particular, when at a switching time a current cell becomes Vac, the head must already be on another (current) cell.
- The length of any noise-free dwell period in which the head is staying on clean positions is at most T .

We say that a cell *dies* in a history if it becomes Vac. Let

$$\text{Histories}_M$$

denote the set of all possible histories of M .

The above definition can be *localized* to define a history over a *space-time rectangle* $I \times J$, where it is always understood that $h \in I$ for all times $t \in J$, that is I contains the head throughout the time interval considered. \lrcorner

In Section 6.4 below, we will define a certain subset of possible histories, called *trajectories*. The set of trajectories of M will be denoted by

$$\text{Trajectories}_M.$$

The definition of generalized Turing machines will be finished only by the definition of the set of trajectories. But in order to motivate their choice, we first introduce the notions of simulation and a hierarchy of simulations.

6.2 Simulation

Until this moment, we used the term “simulation” informally, to denote a correspondence between configurations of two machines which remains preserved during the computation. In the formal definition, this correspondence will essentially be a code $\varphi = (\varphi_*, \varphi^*)$. The *decoding* part of the code is the more important. We want to say that machine M_1 simulates machine M_2 via simulation φ if whenever (η, Noise) is a trajectory of M_1 then (η^*, Noise^*) , defined by $\eta^*(\cdot, t) = \varphi^*(\eta(\cdot, t))$, is a trajectory of M_2 . Here, Noise^* is computed by an appropriate mapping.

We will make, however, two refinements. First, we may weaken the condition by requiring it only for those η for which the initial configuration $\eta(\cdot, 0)$ has been obtained by encoding, that is it has the form $\eta(\cdot, 0) = \varphi_*(\xi)$. This gives a role to the encoding function in the definition.

But there is a more complex refinement. When a colony is in transition between encoding one simulated value to encoding another one, there may be times when the value represented by it before the transition is already not decodable from it, and the value after the transition is not yet decodable from it. Our way of dealing with these situations is to define the simulation decoding as a mapping Φ^* between *histories*, not just configurations. This allows a certain amount of “looking back”: the map Φ^* can depend on the configurations at the beginning of the “work period”.

It is the mapping Φ^* that will also define Noise^* . A history was defined above in Definition 6.5 as a pair (η, Noise) , so we will have $\Phi^*(\eta, \text{Noise}) = (\eta^*, \text{Noise}^*)$. The set Noise^* will be defined just as in Definition 5.2 of sparsity: by deleting those small isolated parts of Noise that the error-correcting simulation can deal with.

Definition 6.6 (Simulation) Let M_1, M_2 be two generalized Turing machines, and let

$$\varphi_* : \text{Configs}_{M_2} \rightarrow \text{Configs}_{M_1}$$

be a mapping from configurations of M_2 to those of M_1 , such that it maps starting configurations into starting configurations. We will call such a map a *configuration encoding*. Let

$$\Phi^* : \text{Histories}_{M_1} \rightarrow \text{Histories}_{M_2}$$

be a mapping. The pair (φ_*, Φ^*) is called a *simulation* (of M_2 by M_1) if for every trajectory (η, Noise) with initial configuration $\eta(\cdot, 0) = \varphi_*(\xi)$, the history $(\eta^*, \text{Noise}^*) = \Phi^*(\eta, \text{Noise})$ is a trajectory of machine M_2 .

We say that M_1 *simulates* M_2 if there is a simulation (φ_*, Φ^*) of M_2 by M_1 . \lrcorner

6.3 Hierarchical codes

Recall the notion of a code in Definition 1.2.

Definition 6.7 (Code on configurations) Consider two generalized Turing machines M_1, M_2 with the corresponding alphabets and transition functions, and an integer $Q \geq 1$. We require $B_2 = QB_1$. Assume that a block code $\psi_* : \Sigma_2 \rightarrow \Sigma_1^Q$ is given, with an appropriate decoding function, ψ^* . Symbol $a \in \Sigma_2$, is interpreted as the content of some tape square.

This block code gives rise to a *code on configurations*, that is a pair of functions

$$\varphi_* : \text{Confs}_{M_2} \rightarrow \text{Confs}_{M_1}, \quad \varphi^* : \text{Confs}_{M_1} \rightarrow \text{Confs}_{M_2}$$

that encodes some (initial) configurations ξ of M_2 into configurations of M_1 . Let ξ be a configuration of M_2 with $\xi.\text{cur-cell}_j = \xi.\text{pos} + jB_2$. We set $\varphi_*(\xi).\text{pos} = \xi.\text{pos}$, $\varphi_*. \text{cur-cell}_j = \varphi_*(\xi).\text{pos} + jB_2$, and

$$\varphi_*(\xi).\text{tape}(iB_2, \dots, (i+1)B_2 - B_1) = \psi_*(\xi.\text{tape}(i)).$$

A configuration ξ is called a *code configuration* if it has the form $\xi = \varphi_*(\zeta)$. \lrcorner

Definition 6.8 (Hierarchical code) For $k \geq 1$, let Σ_k be an alphabet, of a generalized Turing machine M_k . Let $Q_k > 0$ be an integer colony size, let φ_k be a code on configurations defined by a block code

$$\psi_k : \Sigma_{k+1} \rightarrow \Sigma_k^{Q_k}$$

as in Definition 6.7. The sequence (Σ_k, φ_k) , $(k \geq 1)$, is called a *hierarchical code*. For the given hierarchical code, the configuration ξ^1 of M_1 is called a *hierarchical code configuration* if a sequence of configurations ξ^2, ξ^3, \dots of M_2, M_3, \dots exists with

$$\xi^k = \varphi_{*k}(\xi^{k+1})$$

for all k . (Of course, then whole sequence is determined by ξ^1 .)

Let M_1, M_2, \dots be a sequence of generalized Turing machines, let $\varphi_1, \varphi_2, \dots$ be a hierarchical code for this sequence, let ξ^1 be a hierarchical code configuration for it, where ξ^k is an initial configuration of M_k for each k . Let further be given a sequence of mappings $\Phi_1^*, \Phi_2^*, \dots$ such that for each k , the pair (φ_{*k}, Φ_k^*) , is a simulation of M_{k+1} by M_k . Such an object is called a *tower*. \lrcorner

The main task of the work will be the construction of a certain tower. Since the simulation property is highly nontrivial, so is the existence of towers. Generalized Turing machines with the particular form of the notion of a trajectory are introduced since they can be proved to form towers.

6.4 Trajectories

This subsection completes the definition of a generalized Turing machine. The set of trajectories Trajectories_M is defined in terms of constraints imposed on the burst-free parts of a history. We discuss these properties first informally.

Transition Function This property says (in more precise terms) that in the absence of noise and in a clean area, the transition function is obeyed.

The Spill Bound limits the extent to which a disordered interval can spread while the head is in it.

Escape limits the time that the head can spend in a small area.

Attack Cleaning helps erode the disorder as the head repeatedly enters and leaves it.

Pass Cleaning limits the number of back-and-forth slides that can happen over a large disordered area (under certain conditions). Otherwise every time the head reaches, say, its left end, a new burst could expand the disorder. As these bursts are not visible in the simulated trajectory η^* , the Spill Bound property of η^* could be violated.

Definition 6.9 Suppose that at times t' before a switching time t but after any previous switch, the current cell-pair has state \mathbf{a} , further after the switch the same cell pair has state \mathbf{a}' (including when one of these cells dies).

We say that the switch is *dictated by the transition function* if

$$(\mathbf{a}', \text{sign}(\hat{h}_0(t') - \hat{h}_0(t))) = \tau(\mathbf{a}),$$

further $a'_0 \cdot \text{Adj} = 1$ if and only if $\hat{h}_0(t') - \hat{h}_0(t) \in \{-B, 0, B\}$. ┘

We will need the following constants, to be fixed later.

$$c_{\text{attack}} < \beta/2, \quad c_{\text{spill}}, \quad 5 < \theta < \beta, \quad c_{\text{pass}} \leq \theta/6, \quad c_{\text{esc}}. \quad (6.4)$$

Peter remark 1. Some details of the following definition need to be adjusted yet. ┘

Definition 6.10 (Trajectory) A history (η, Noise) of a generalized Turing machine (6.1) with $\eta(t) = (A(t), h(t), \hat{\mathbf{h}}(t))$ is called a *trajectory* of M if the following conditions hold, in any noise-free space-time interval $I \times J$.

Transition Function Suppose that there is a switch, and the current cell-pair is inside the clean area, by a distance of at least $2.5B$. Then the new state of the current cell-pair, and the direction towards the new current head position are dictated by the transition function. If a new current cell did not exist before then it is adjacent to one of the previous ones before the switch. The only change on the tape occurs on the interval enclosing the new and old current cells. Further, the length of the dwell period is bounded by T .

Spill Bound An interval of disorder can spread by at most $c_{\text{spill}}B$ on either side while it contains the head.

Escape The head will leave any interval of size $\leq \lambda B$ with $1 \leq \lambda \leq 3\beta$ within time $c_{\text{esc}}\lambda^2 T$.

Attack cleaning Suppose that the current cell-pair is at the end of a clean interval of size $\geq (c_{\text{attack}} + 1)B$. Suppose further that the transition function directs the head right. Then by the time the head comes back to $x - c_{\text{attack}}B$, the right end of the clean interval containing it advances to the right by at least B .

A similar property is required when “left” and “right” are interchanged.

Pass Cleaning This property will be defined below in Definition 6.13. It is motivated by Example 2.5.

┘

Our construction will set $\pi = O(k)$ for a machine on level k of the simulation, that is $O(k)$ passes will restore organization to level k . (Essentially, at most three slides are needed to raise the level.) But some restriction is needed: if between the noiseless passes there is a large number of *intrusions* that are not passes, then each of these intrusions can bring a little new disorder, possibly undoing the effect of passes. Limiting the *number* of intrusions does not seem feasible, since there can be too many local intrusions. Instead, we will limit the number of a certain kind of rectangles covering them. This will help since in our noise model, if a space-time area is noiseless for the simulated machine M^* then essentially every $2\gamma \cdot (QB) \times (UT)$ rectangle contains at most one burst (of size $\beta \cdot B \times T$).

Definition 6.11 For a space-time path P of the head and an interval I , let $P_{u,v}$ be its part between times u and v . We write $P_{\cdot,v}$ for the part from start to v and $P_{u,\cdot}$ for the part from u to end.

A part $P_{u,v}$ is a *pass* of I if P enters I at time u and leaves it on the other side at time v (without leaving I before). \lrcorner

Definition 6.12 (*W-segments*) Consider a path P over some time interval $J = [t, u)$ with the property that no burst occurs on it while it is in space interval I (bursts might occur on it while it is outside I). Break it up into segments in such a way that each segment has space projection $\leq W$ and only the last segment can have space projection $< W$. Fix such a segmentation and call it the *W-segmentation* of P ; its parts are called the *W-segments*. In the *W-segmentation* of P we call the *W-segments* intersecting I but not belonging to a pass of I the *W-intrusions* of P into I . \lrcorner

Note that if P is noiseless and $W = \lambda B$ with $1 \leq \lambda \leq 3\beta$ then by the Escape property, the time spent in each *W-segment* of P is at most $c_{\text{esc}}\lambda^2 T$.

Recall the constants defined in (6.4).

Definition 6.13 (*Pass cleaning property*) Suppose that a path P makes at least π pairs of passes over interval I of size $\leq \theta B$, with at most a number π^2 of θB -intrusions. Then at some time during P the interior $\text{Int}(I, c_{\text{pass}}B)$ of I becomes clean. \lrcorner

The trajectory properties can clearly be localized to some space-time rectangle just as the definition of history was.

7 Simulation structure

In what follows we will describe the program of the reliable Turing machine (more precisely, a simulation of each M_{k+1} by M_k as defined above). Most of the time, we will refer to M_k as M and to M_{k+1} as M^* . Cells will be grouped into colonies, where $Q = B^*/B$ is the colony size. The behavior of the head on each colony simulates the head of M^* on the corresponding cell of M^* . The process takes a number of steps, constituting a *work period*.

Definition 7.1 We will have a constant

$$c_{\text{sim}} \tag{7.1}$$

with the property that the maximum number of steps that any simulation work period is at most $c_{\text{sim}}\beta Q^2/2$. Let

$$U = c_{\text{sim}}\beta Q^2, \quad T^* = UT.$$

\lrcorner

Machine M will perform the simulation even if the noise in which it operates is $(\beta(B, T), \gamma(B^*, T^*))$ -sparse. By the above definitions, sparsity means that noise comes in *bursts* that are confined to rectangles of size β (affecting at most β consecutive tape cells), and are separated from each other in time in such a way that there is at most one burst in any γ neighboring work periods. A design goal for the program is the following:

Goal 7.2 (Local correction) *Correct a burst within space and time comparable to its size, and much smaller than the size of a simulation work period.*

The program is rather complex, but there are some limitations to a modular presentation in which after the presentation of each part it is shown what properties of the simulation it guarantees. Indeed, noise occurring during any part of the program can change the environment of the head into some state corresponding to an arbitrary other part.

Modes were introduced in Section 4.2. Ordinary simulation proceeds in the normal mode. To see whether the basic structure supporting this process is broken somewhere, each step will check whether the current cell-pair is *coordinated* (see Definition 8.5). If not then the state of the current pair will be changed into the *healing* mode. We will also say that *alarm* will be called. On the other hand, the state may enter into *rebuilding* mode on some indications that healing fails.

7.1 Head movement

The global structure of a work period is this:

Computation phase The new state of the simulated cell pair, and the simulated direction (called the drift) is computed. Then the “meaningfulness” of the result is checked. During this phase the head sweeps, roughly, back-and-forth between the ends of the base colony-pair.

Transfer phase The head moves into the neighbor colony in the simulated head direction called drift (building or rebuilding a bridge if needed). Redundancy is used to protect this movement from a burst.

The number of the current sweeps within a work period is contained by a field *Sweep*.

Definition 7.3 (Front) During normal computation, as the head sweeps in a certain direction, it increases *Sweep* field by 1. Let us call the last site in which this increase occurred, the *front*. ┘

Globally in a configuration, due to disorder, there may be more than one front, but locally we can talk about “the” front without fear of confusion. We

can read off the direction of the sweep s from its parity:

$$\text{dir}(s) = (-1)^{s+1}. \quad (7.2)$$

Due to the zigging and feathering requirements, in Section 2.4, there will be two kinds of turns: *small* turns and *big* turns. The process uses the constants

$$c_{\text{turn-limit-1}}, c_{\text{turn-limit-2}}, f = 5 \quad (7.3)$$

to be fixed later, and the parameters,

$$Z = \pi^{2+\varepsilon} \quad F = Z\pi^{2+\varepsilon} \quad (7.4)$$

with $\varepsilon > 0$ to be fixed later, which allows us to define the earlier used parameter

$$\text{PadLen} = F \log Q. \quad (7.5)$$

The choices will be motivated in Sections 7.1.2 and 12.3. When the simulation needs the head to turn, this will in general not happen immediately. Turns will be handled by a field

$$\text{Turned} \in \{0, 1, 2\}.$$

The value 1 shows the memory of a recent turn. The value 2 shows the memory of a recent big turn within distance F . When a cell is passed in which $\text{Turned} > 0$ was found then this value is set to 0.

Small turn Move until you found a cell with $\text{Turned} \neq 1$, then turn. If you did not find it in $c_{\text{turn-limit-1}}$ steps then call alarm. Set $\text{Turned} \leftarrow 1$ in this cell.

Small turns will be done during zigging and healing (see later).

Big turns Move until you found F consecutive cells with $\text{Turned} < 2$ the last one of which has $\text{Turned} = 0$, then make a turn. On the cell after the turn, set $\text{Turned} \leftarrow 1$, and also on the first interval of F cells after the turn, set $\text{Turned} \leftarrow 2$. If you did not find it in $c_{\text{turn-limit-2}}Q$ steps then call alarm: such an event will be called a *turn-starvation*. The control can be accomplished by a counter, say called *FeatherDepth*.

Big turns will be done during the ends of sweeps in simulation and rebuilding.

7.1.1 Zigging

A burst could turn the head back in the middle of its sweep, as pointed out in Example 2.2. To detect such an event promptly (in accordance with Goal 7.2), a *zigzag* movement will be superimposed on the sweeping, so that the head can check the consistency of a few cells around the front. (Their number is not constant, it will depend on π .) The process creates a *frontier zone* of about $2Z$ cells in both directions around the front, where Z was defined in (7.4). The field

$$\text{ZigAddr} \in [-2Z, 2Z).$$

of the current cell shows the cell's address position in the frontier zone. The zone itself is the interval of cells in which *ZigAddr* is defined. After every f steps of progress, the head will perform a forward-backward-forward zigzag checking the zone. For this it uses *small turns* defined above, so it may need a few more steps to find a turning point. Due to the spacing of all turns, under normal conditions, this will need no more than f steps (as will be proved).

On backward zig (with respect to the sweep direction), we expect the zig addresses to grow continuously from $-2Z$ to $2Z$; the head does not change them. On forward zig, the addresses will look like $\dots, i-1, i, i+j, i+j+1, \dots$, for some $0 < j \leq f$. If $\text{SwDir} > 0$ then $i+j$ will be changed to $i+1$, otherwise i will be changed to $i+j-1$.

- Remarks 7.4**
1. The need for backward zigging was explained in Example 2.2.
 2. The forward-zigging property and the frontier zone will be used in the poof of Lemma 12.14. We will need to show there that in the absence of new noise, after a constant number of passes, a clean interval J of size, say, $\geq 6ZB$ will extend to the right until it reaches the end of a colony pair or of a rebuilding area. Forward zigging helps prevent the disorder from turning the head back prematurely.
 3. The choice $Z = \pi^\epsilon$ will be justified in the same proof: it lower-bounds the size of new noise capable of turning back the head.

┘

7.1.2 Feathering

Here we motivate the definition of *small turns* and *big turns* introduced above. Example 2.3 above pointed out the need for the following feature of the simulation program.

Definition 7.5 (Feathering) A Turing machine execution is said to have the *c-feathering* property if the following holds. If the head turned back at a position x at some time, then next time it comes within distance c of x , it can turn back only at least c steps beyond x . ┘

(The name “feathering” refers to the picture of the path of the head in a space-time diagram.) The following example suggests that any computation can be reorganized to accomodate feathering, without too much extra cost.

Example 7.6 (1-feathering) Suppose that, arriving from the left at position 1, the head decides to turn left again. In repeated instances, it can then turn back at the following sequence of positions:

1, 2, 1, 3, 1, 2, 1, 4, 1, 2, 1, 3, 1, 2, 1, 5, 1, 2, 1, 3, 1, 2, 1, 4, 1, 2, 1, 3, 1, 2, ...

┘

If in the original execution the head turned back t consecutive times to the left from position p , then now it will turn back from somewhere in a zone of size $O(\log t)$ to the right of p in each of these times. (Computing the exact turning point is not necessary.)

Normally, a small turn attempt would take longer than $O(1)$ steps and a big turn attempt would take longer than $O(F \log Q)$ steps only in the presence of large disorder (those cases will be analyzed in the context of large disorder). We give here only an informal argument justifying this statement; a formal proof must wait until a complete definition of the simulation and of the notion of large disorder.

For the moment, let us call cells *red* with by $Turned = 1$ and *orange* with $Turned = 2$. Red cells can delay small turns; they can only arise in the forward turn of a zig in the previous sweep. Zigs are by the definition spaced by $\geq f$ cells apart. Big turns are delayed by orange cells; they are also the ones to create orange cells. But they will be organized as in Example 7.6, therefore a big turn attempt will be delayed by at most F times the logarithm of the total number of big turns inside a colony or a rebuild area.

The turns at the end a rebuilding area are more problematic since there may be nothing known about the cells over which it is extended—so it may indeed happen that a place to turn is not reached before alarm is called: this is what we called *turn starvation* above.

The simulated Turing machine will also have the feathering property, therefore the simulation will not turn back from one and the same colony repeatedly, without having passed it in the meantime.

Remark 7.7 The size of the parameter F is motivated by the proof of Lemma 12.14. Here is a sketch of the argument (it can be safely skipped now). At some time t_0 in some interval I we will have clean subintervals $J_k(t_0)$, $k = 1, 2, \dots$ of size about $6ZB$ in which no burst will appear, and which are separated from each other by areas of size $O(\pi^2 ZB) \ll FB$. For times $t > t_0$ we will consider the maximal clean intervals $J_k(t)$ containing the middle of $J_k(t_0)$.

Assume that the head passes over I noiselessly left to right and later from right to left with no bursts. If the head moves in a zigging way to the right then the attack property will clean out the area between $J_i(t)$ and $J_{i+1}(t)$, joining them. This does not happen only in case of a programmed turn from the right end of $J_i(t)$, in the course of the simulation or rebuilding. But then in the next pair of passes over I , the feathering property implies that the programmed turn from the end of $J_i(t)$ is at least a distance FB to the right. Our choice of F implies that then $J_i(t)$ will be joined to $J_{i+1}(t)$. So two noise-free passes would join all the intervals $J_i(t)$ into a clean area. \lrcorner

We introduce a convenient notation.

Notation 7.8 Suppose that within a procedure a turn attempt is made after a move of n cells in some direction. Then the head may not be able to turn right away, the turn attempt moves it a few more steps in the same direction. In this case we will say that the head moves n^+ cells and turns. Thus n^+ is greater than n , but by not much: only an amount $O(1)$ for small turns, and $O(F \log Q)$ for big turns. \lrcorner

7.2 Computation phase

The first phase of the simulation, called the Compute rule, computes new values for current cell-pair of the simulated machine M^* represented by the current (base) colony-pair, and the direction of the move of the head of M^* . During the execution of this rule, the head sweeps the base colony-pair. The cell state of M^* will be stored on the track *Info* of the representing colony. The move direction of M^* will be represented in the *Drift* field of *each* cell of the base colony-pair (so the whole track must be filled with -1 's or 1 's).

Recall Definition 5.4. The rule Compute will rely on a certain fixed $(\beta, 3)$ burst-error-correcting code, moreover it expects that each of the words found on the *Info* track is 2-compliant. The rule ComplianceCheck checks whether a word is 2-compliant.

The rule Compute essentially repeats 3 times the following *stages*: decoding, applying the transition, encoding. Then it calls ComplianceCheck; if the latter

fails it will mark the colony for rebuilding. It uses some additional tracks. A track called

Work

can be used for auxiliary computation. A useful track is called

Adj'.

At the time of the computation, on this track of the left base colony, each cell should contain a single symbol in $\{0, 1\}$, according to whether the previous left base colony was adjacent or not. In more detail:

1. For $j = 1, \dots, 3$ do
 - a. Calling by \mathbf{a} the pair of strings found on the *Info* track of the interior of the base colonies, decode it into the pair of strings $\tilde{\mathbf{a}} = v^*(\mathbf{a})$ (this should be the current state of the simulated cell-pair), and store it on some auxiliary track in the base colony-pair. Do this by simulating the universal machine: $\tilde{\mathbf{a}} = \text{Univ}(p_{\text{decode}}, \mathbf{a})$. Then replace $\tilde{a}_0.\text{Adj}$ with the value found on the *Adj'* track.
 - b. Compute $(\mathbf{a}', d) = \tau^*(\tilde{\mathbf{a}})$. Since the program of the transition function τ^* is not written explicitly anywhere, this “self-simulation” step is elaborated in Section 7.3.
 - c. Write the encoded new cell states $v_*(\mathbf{a}')$ onto the *Hold[j].Info* track of the interior of the base colony-pair. Write d into the *Hold[j].Drift* field of each cell of the left base colony.
 If the new state a'_i is a vacant one for $i = 0$ or 1 , that is $a'_i.\text{Kind}^* = \text{Vac}^*$, then write 1 onto the *Hold[j].Doomed* track of the corresponding (left or right) base colony—else write 0.
2. Sweeping through the base colony pair, at each cell compute the majority of *Hold[j].Info*, $j = 1, \dots, 3$, and write it into the field *Info*. Proceed similarly, and simultaneously, with *Drift*.
3. For $j = 1, \dots, 3$, call *ComplianceCheck* on the *Info* track, and write the resulting bit into the *Compliant_j* track.
 Then pass through the colony and turn each cell in which the majority of *Compliant_j*, $j = 1, \dots, 3$ is false, into a stem cell (thus destroying the colony if the result was false everywhere).

It can be arranged—and we assume so—that the total number of sweeps of this phase, and thus the starting sweep number of the next phase, depends only on Q .

7.3 Forced self-simulation

Here we elaborate step 1b of Section 7.2.

7.3.1 New primitives

We will make use of the tracks *Work* mentioned above, and the track

Index

that can store a certain address of a colony.

Recall from Section 4.2 that the program of our machine is a list of nested “**if condition then instruction else instruction**” statements. As such, it can be represented as a binary string

R .

If one writes out all details of the construction of the present paper, this string R becomes explicit, an absolute constant. But in the reasoning below, we treat it as a parameter.

Let us provide a couple of *extra primitives* to the rules. First, they have access to the parameter k of machine $M = M_k$, to define the transition function

$\tau_{R,k}(\mathbf{a})$.

The other, more important, new primitive is a special instruction

WriteProgramBit

in the rules. When called, this instruction makes the assignment $Work \leftarrow R(Index)$. This is the key to self-simulation: *the program has access to its own bits*. If $Index = i$ then it writes $R(i)$ onto the current position of the *Work* track.

7.3.2 Simulating the rules

By convention, in our fixed flexible universal machine Univ, program p and input x produce an output $Univ(p, x)$. The structure of all rules is simple enough that they can be read and interpreted by Univ in reasonable time:

Theorem 2 *There is a constant string called Interpr with the property that for all positive integers k , string R that is a sequence of rules, and a pair of bit strings $\mathbf{a} = (a_0, a_1)$ with $a_j \in \Sigma_k$,*

$$Univ(Interpr, R, 0^k, \mathbf{a}) = \tau_{R,k}(\mathbf{a}).$$

The computation on Univ takes time $O(|R| \cdot |\mathbf{a}|)$.

The proof parses and implements the rules in the string R ; each of these rules checks and writes a constant number of fields.

Implementing the `WriteProgramBit` instruction is straightforward: Machine Univ determines the number i represented by the simulated *Index* field, looks up $R(i)$ in R , and writes it into the simulated *Work* field.

There is no circularity in these definitions:

- The instruction `WriteProgramBit` is written *literally* in R in the appropriate place, as “`WriteProgramBit`”. The string R is *not part* of the rules (that is of itself).
- On the other hand, the computation in $\text{Univ}(\text{Interpr}, R, 0^k, \mathbf{a})$ has *explicit* access to the string R as one of the inputs.

Let us show the computation step invoking the “self-simulation” in detail. In the earlier outline, step 1b of Section 7.2 said to compute $\tau^*(\tilde{\mathbf{a}})$ (for the present discussion, we will just consider computing $\tau^*(\mathbf{a}) = \tau_{k+1}(\mathbf{a})$), where $\tau = \tau_k$, and it is assumed that \mathbf{a} is available on an appropriate auxiliary track. We give more detail now of how to implement this step:

1. Onto the *Work* track, write the string R . To do this, for *Index* running from 1 to $|R|$, execute the instruction `WriteProgramBit` and move right. Now, on the *Work* track, replace it with $\langle \text{Interpr}, 0^{k+1}, R, \mathbf{a} \rangle$. Here, string *Interpr* is a constant, so it is just hardwired. String R already has been made available. String 0^{k+1} can be written since the parameter k is available. String \mathbf{a} is available on the track where it is stored.
2. Simulate the universal automaton Univ on track *Work*: it computes $\tau_{R,k+1}(\mathbf{a}) = \text{Univ}(\text{Interpr}, R, 0^{k+1}, \mathbf{a})$ as needed.

This implements the forced self-simulation. Note what we achieved:

- On level 1, the transition function $\tau_{R,1}(\mathbf{a})$ is defined completely when the rule string R is given. It has the forced simulation property by definition, and string R is “*hard-wired*” into it in the following way. If $(\mathbf{a}', d) = \tau_{R,1}(\mathbf{a})$, then

$$a'_0.\text{Work} = R(a_0.\text{Index})$$

whenever $a_0.\text{Index}$ represents a number between 1 and $|R|$, and the values $a_0.\text{Sweep}$, $a_0.\text{Addr}$ satisfy the conditions under which the instruction `WriteProgramBit` is called in the rules (written in R).

- The forced simulation property of the *simulated* transition function $\tau_{R,k+1}(\cdot)$ is achieved by the above defined computation step—which *relies on* the forced simulation property of $\tau_{R,k}(\cdot)$.

Remark 7.9 This construction resembles the proof of Kleene’s fixed-point theorem, and even more some self-reproducing programs (like a program p in the language C causing the computer to write out the string p). \lrcorner

7.4 Transfer phase

In the transfer phase the Transfer rule takes over: control will be transferred to the neighbor colony in the direction of the simulated head movement: this is called the direction of the transfer, or the *drift*. As we will see the transfer phase takes only a constant number of sweeps: three pairs of sweeps should be sufficient (see the details below). During this phase, the range of the head includes the base colony pair and a neighbor colony in the direction of the drift called *target colony*, including possible bridges between them. The bridges present some extra complication—let us address it.

The sweep number in which we start transferring in direction δ is called $\text{TransferSw}(\delta)$, the *transfer sweep*. We set $\text{TransferSw}(-1) = \text{TransferSw}(1) + 1$.

Definition 7.10 (Adjacency of cells) Cells a and b are *adjacent* if $|a - b| = B$. Otherwise, if $B < |a - b| < 2B$, then a and b are two *non-adjacent neighbor cells*. For the sake of the present discussion, a *colony* is a sequence of Q adjacent cells whose *Addr* value runs from 0 to $Q - 1$. It may be *extended* by a bridge of up to $Q - 1$ adjacent cells.

If the bodies of two cells are not adjacent, but are at a distance $< B$ then the space between them is called a *small gap*. We also call a small gap such a space between the bodies of two colonies. On the other hand, if the distance of the bodies of two colonies is $> B$ but $< QB$ then the space between them is called a *large gap*. \lrcorner

Before the transfer phase, the base colony pair consists of two colonies, having cells of kind Member_0 and Member_1 . There may be a bridge between them, of $< Q$ cells of kind Bridge adjacent to each other, and extending one of the base colonies (see Definition 4.6 and the remark following it).

The transfer phase moves the base colony pair left if $\text{Drift} = -1$, and right, if $\text{Drift} = 1$ in the left base colony. Consider $\text{Drift} = -1$ first. The kind of cells of the left base colony will turn from Member_0 to Member_1 . Then a bridge is built to the left (extending the -1 value on the *Drift* track). This bridge can override an opposite old bridge (“old” meaning that its *Sweep* is maximal) or move into an empty area, kill other bridge cells or stem cells if they are not adjacent, while it extends. There are two ways that this bridge extension can finish.

- It stops at the boundary of another colony on the left (which consists of cells of type Member_0). (This can happen in the very first step, in which case no bridge is built.) Then this colony becomes the new left base colony, and the head moves to its left end (extending again the -1 value on the *Drift* track).
- The bridge reaches the size of Q cells. In this case, the bridge will be converted into the new left base colony, the kind of its cells becoming Member_0 .

Recall that the *Adj* field determines whether the current cell is adjacent to the cell where the head came from. After the transfer stage, we write the *Adj'* track of the new left base colony: it becomes 0 if either there is a nonempty bridge, or there is a gap (found with the help of the *Adj* field) between the old and new left base colony. This is done again three times, storing candidate values into $\text{Hold}[j].\text{Adj}'$, repeating and doing a majority vote.

In the last sweep, in the old base colony pair, if the majority of $\text{Hold}[j].\text{Doomed}$, $j = 1, \dots, 3$, is 1 then turn the scanned cell into a stem cell: in other words, carry out the destruction.

The case $\text{Drift} = 1$ is analogous, with the following logical changes. First the bridge between the base colonies is possibly rebuilt, to extend the left base colony. Then the kind of cells of the right base colony will turn from Member_1 to Member_0 , and it becomes the new left base colony. Then a bridge is built to its right (extending the 1 value on the *Drift* track). It stops at the boundary of another colony on the right (which consists of cells of type Member_1), or becomes itself a colony with elements having kind Member_1 .

8 Health

The main part of the simulation uses an error-correcting code to protect information stored on the *Info* track. However, faults can ruin the simulation structure and disrupt the simulation itself. The error-correcting capabilities of the code will preserve the content of the *Info* track as long as the coding-decoding process implemented in the simulation is carried out. The correctness of the process itself relies only on the structural integrity of a configuration; this in turn is maintained with the help of a small number of fields. Below we outline the necessary relations among them allowing the identification and correction of local damage.

A configuration with local structural integrity will be called *healthy*.

Remark 8.1 In all discussions of the health of a configuration $\xi = (A, h, \hat{\mathbf{h}})$, we can ignore the current cell position $\hat{h}_0 = \xi.\text{cur-cell}$, since at all switching times it agrees with $h = \xi.\text{pos}$. \lrcorner

All configurations are made up of homogenous domains (interval).

Definition 8.2 (Homogenous domain) The tuple of fields

$$Core = (Kind, Adj', Drift, Addr, Sweep, ZigAddr)$$

is called the *core*. An interval of non-stem adjacent neighbor cells is a *homogenous domain* if its core variables with the exception of *Addr* and *ZigAddr* have the same value, *Addr* increases left to right, and *ZigAddr* values are either all undefined or are increasing left to right. The *left end* of a domain is the left edge of its first cell, and its *right end* is the right edge of its last cell. ┘

A key property of our definition of health is that it is completely determined by the types of domains allowed and the types of boundaries allowed between them. This enables the local detection of failures of health. The big picture is that cells of a healthy configuration are grouped into gapless colonies, with certain transitional intervals inside and between them.

- There is a base colony-pair, with possibly a bridge between them going in the direction of its own drift. If the sweep inside this colony-pair shows that the Compute rule is being performed then the sweep boundaries allowed inside the colony-pair are the front, and the edges according to past turns at the end of the interior (as in Definition 5.6) of the base colony-pair. Outside each edge is a domain with a lower sweep number.
- During the first sweep of transfer, the base colony (of the pair) in the direction of the drift is possibly extended by a bridge towards the target colony. This bridge may be in the process of consuming an earlier bridge extended in the opposite direction. In the later part of the first sweep of transfer, the bridge already reaches the target colony.

If the bridge extends to a full colony then this is converted to the corresponding kinds of member cells in the second sweep of the transfer.

In the later sweeps of the transfer, domains ahead and behind the front show the changes done by the transfer phase, including the change of *Adj'* in the last sweep.

- Non-base colonies are called *outer colonies*. If an outer colony is not adjacent to its neighbor colony closer to the base, then it is extended by a bridge that covers this gap.

Peter remark 2. Pictures! ┘

Let us define outer cells and boundary types in more detail.

Definition 8.3 (Outer cells, desert) Recall the definition of the sweep value $\text{Last}(\delta)$ from (4.1). For $\delta \in \{-1, 1\}$, if a cell is stem or has $\text{Drift} = \delta$, $\text{Sweep} = \text{Last}(\delta)$ then it will be called a *right outer cell* if $\delta = -1$, a *left outer cell* if $\delta = 1$. ┘

According to this definition, a stem cell is both a left and a right outer cell. Recall Definition 7.3 of the front.

Definition 8.4 (Boundaries) A boundary of a homogenous domain is called *rigid* if its address is the end address of a colony in the same direction. A *boundary pair* is a right boundary followed by a left boundary at distance $< B$. It is a *hole* if the distance is positive. It is *rigid* if at least one of its elements is. ┘

The health of a configuration ξ of a generalized Turing machine M is defined over a certain interval A . It depends on $\xi_{\text{tape}}(A)$, further on whether ξ_{pos} is in A and if it is, where. But we will mention the interval A explicitly only where it is necessary. In particular, some of the structures described above may fall partly or fully outside A .

Defining health formally can be done mechanically on the basis of the above informal description, but the detailed description would be tedious.

A tape configuration is called *healthy* on an interval A when there is a head position (possibly outside A) that turns it into a healthy configuration.

Health only depends on the *Core* track and the lack of heal or rebuild marks on the tape. Note that in a healthy configuration every cell's *Core* field determines the direction in which the front is found, from the point of view of the cell.

A violation of the health requirements can sometimes be noted immediately:

Definition 8.5 (Coordination) The current cell pair is *coordinated* if it is possible for them to be together in a healthy configuration. ┘

In normal mode, if lack of coordination is discovered then the healing procedure is called. The following lemmas show how local consistency checking will help achieve longer-range consistency.

Lemma 8.6 *In a healthy configuration, each Core value along with Z determines uniquely the Core value of the other cell of the current cell pair, with the following exceptions.*

- *During the first transferring sweep, while creating a bridge, the front can be a stem cell or the first cell of an outer colony.*
- *Every jump backward from a target colony can land on the last cell of a bridge (whose address is not recorded in the cell from which the head is jumping) or the last cell of the base colony.*

Proof. To compute the values in question, use Z to find the address at the front, and refer to the properties listed above.

Peter remark 3. Elaborate! ┘

□

Lemma 8.7 (Health extension) *Let ξ be a tape configuration that is healthy on intervals A_1, A_2 where $A_1 \cap A_2$ contains a whole cell body of ξ . Then ξ is also healthy on $A_1 \cup A_2$.*

Proof. The statement follows easily from the definitions.

Peter remark 4. Elaborate! ┘

□

In a healthy configuration, the possibilities of finding non-adjacent neighbor cells are limited.

Lemma 8.8 *An interval of size $< Q$ over which the configuration ξ is healthy contains at most two maximal sequences of adjacent non-stem neighbor cells.*

Proof. Indeed, by definition a healthy configuration consists of full extended colonies, with possibly stem cells between them. An interval of size $< Q$ contains sequences of adjacent cells from at most two such extended colonies. □

Lemma 8.9 *In a healthy configuration, a cell's state shows whether it is an outer cell, colony cell, bridge cell or target colony cell. It also shows whether the cell is to the left or to the right of the front.*

The Core track of a homogenous domain can be reconstructed from any of its cells.

Proof. The information mentioned in the first sentence is explicit in some fields. About the front: if the cell is outer then *Drift* shows its direction from the front. In case it is not outer, then the parity of *Sweep* shows it: odd values are on the left, even ones on the right.

The reconstruction is a direct consequence of the definitions. □

9 Healing and rebuilding

9.1 Annotation, admissibility, stitching

In this section we show how to correct configurations of machine M that are “almost” healthy.

Definition 9.1 (Annotation) An *annotated configuration* is a tuple

$$(\xi, \chi, \mathcal{I}),$$

with the following meaning.

ξ is a configuration.

\mathcal{I} is a set of disjoint intervals called *islands*. Their complement is clean.

χ is a healthy configuration differing from ξ only in the islands.

The *base colony* of (ξ, χ, \mathcal{I}) is that of χ . The head is *free* when it is not in any island, and the observed cell is in normal mode, coordinated with its pair. \lrcorner

Definition 9.2 (Admissibility) An annotation is *admissible* on an interval K if the following holds on any subinterval J of K of size QB :

- a1) The disorder of J is covered by 3 intervals, each of size $\leq \beta B$.
- a2) There are at most 3 islands in J , each of size at most $c_{\text{island}}\beta B$.

A configuration is *admissible* on K if it can be annotated in a way admissible on K . \lrcorner

Let us argue (informally, without replacing any later, formal argument), that local correction does not have to deal with more than three islands in any area of size QB .

Example 9.3 (Three islands) Suppose that the head has arrived at a colony C from the left, performs a work period and then passes to the right. In this case, if no new noise burst occurs then we expect that all islands found in C will be eliminated by the healing procedure. On the other hand, a new island I_1 can be deposited (in the last pass). We can assume that there is no island on the left of I_1 within distance QB , since the noise burst causing it would have been too close to the noise burst causing I_1 .

Consider the next time (possibly much later), when the head arrives (from the right). If it later continues to the left, then the situation is similar to the above. Island I_1 will be eliminated, but a new one may be deposited. But what if the head turns back right at the end of the work period? If I_1 is close to the left end of C , then due to the feathering construction, the head may never reach

it to eliminate it; moreover, it may add a new island I_2 on the right of I_1 . We can also assume, similarly to the above, that there is no island on the right of I_2 within distance QB . (In the work period where the head deposits island I_2 near the left colony end, it may repeatedly dip into I_2 at the descending end of a zig at a right turn. During this dip it can expand the islands I_1, I_2 and then emerge on the right, with the healing unfinished.)

When the head returns a third time (possibly much later), from the right, the feathering requirement implies that it will leave on the left. The islands I_1, I_2 will be eliminated as they are passed over but a possible new island I_3 , created by a new burst (before, after or during the elimination), may remain. We can also assume, similarly to the above, that there is no island on the right of I_3 within distance QB . So at least temporarily, it is possible to have three islands in C . \lrcorner

We will show that a configuration admissible over an interval of a certain size can be locally corrected; moreover, in case the configuration is clean (no disorder), this correction can be carried out by the machine M itself. For the time being, we will just talk about an admissible configuration, without specifying the interval K .

Definition 9.4 (Substantial domains) Let $\xi(A)$ be a tape configuration over an interval A . A homogenous domain of size at least $7c_{\text{island}}\beta B$ will be called *substantial*. The area between two neighboring maximal substantial domains or between an end of A and the closest substantial domain in A will be called *ambiguous*. It is *terminal* if it contains an end of A . Let

$$\Delta = 39c_{\text{island}}\beta,$$

$$E \geq 6\Delta.$$

\lrcorner

Lemma 9.5 Consider an admissible configuration. In a substantial domain, each half contains at least one cell outside the islands.

If an interval of size $\leq QB$ of a tape configuration ξ differs from a healthy tape configuration χ in at most three islands, then the size of each ambiguous area is at most ΔB .

Proof. The first statement is immediate from the definition of substantial domains. There are at most 3 boundary pairs in χ at a total size of $3B$, and 3 islands of size $\leq c_{\text{island}}\beta B$. There are at most 5 non-substantial domains of sizes $< 7c_{\text{island}}\beta B$ between these: this adds up to

$$< (3 + 3 \cdot c_{\text{island}}\beta + 5 \cdot 7 \cdot c_{\text{island}}\beta)B < 39 \cdot c_{\text{island}}\beta B = \Delta B.$$

□

The following lemma forms the basis of the healing algorithm.

Lemma 9.6 (Stitching) *In an admissible configuration, inside a clean interval, let U, W be two substantial domains separated by an ambiguous area V . It is possible to change the tape on U, V, W using only information in U, W in such a way that the tape configuration over $U \cup V \cup W$ becomes healthy. Moreover, it is possible for a cellular automaton to do so gradually, keeping admissibility, and changing the tape in U or W or enlarging U or W gradually at the expense of V .*

The machine in question can make its turns via “small turns” as defined in Section 7.1.

Proof. We distinguish several cases, based on the kind of domains involved. At any step, if we find that $U \cup V \cup W$ is healthy then we stop.

1. Assume that U, W both belong to the same extended colony: either an extended base colony, or an outer extended colony, or the target.

In this case, by Lemma 8.9, the content of the *Core* track of V in the healthy configuration is completely determined by that of U, W . So the intervals U and W , which will be recognized as the substantial ones as opposed to V , can be gradually extended towards each other in any order, overtaking V .

If U and W belonged to different sides of the front, then the new front will be the new boundary between U and W .

2. Suppose now that U, W do not belong to the same extended colony, but they are on the same side of the front: without loss of generality, to the left of it. In this case, only U can be outer or a target: suppose that it is outer. Now W is either in a target or in the base colony. The base colony cannot have an extension to the left, because given that the front is to the right, the same sweep that created the extension would have created already a target, separating U (which being outer is not in the target) from W too much. So in both cases, W is close to its colony’s left end, which is in V . It must be extended to this left end. Following this, U must be extended until it reaches W . The *Sweep* values can be propagated without change, there is no need to coordinate them between U and W .

Suppose that U is a target, then W belongs to the extended base colony. Then U must be extended until its endcell, and then W must be extended to meet U . Since both U and W are in the interior, the *Sweep* values must be equal, so they can be extended without change.

3. Suppose that U and W belong to different extended colonies, with the front between them.

The *Sweep* values can be extended without change, there is no need to coordinate them between U and W .

Suppose that U contains no bridge cells: in this case extend it to the right until it hits a colony endcell. If you overlap W , decrease W accordingly. Due to admissibility, only bridge cells of W can be overwritten in this way (colonies don't overlap in a healthy configuration). Similarly, if W contains no bridge cells then extend it to the left, until it hits a colony endcell. Again, only bridge cells of U can be overwritten in this way.

Only one of U and W can be in an outer extended colony, suppose that U is. If W is a target then we already extended it to its limit. Now extend the bridge of U to meet it. If W is in an extended base colony then extend its bridge until either meets U or reaches full length. In the latter case, extend a bridge of U to meet W .

If U is in a target then it is already at its right end: we extend a bridge of W to meet it.

□

9.2 The healing procedure

Structure repair will be consist of two procedures. The first one, called *healing*, performs only local repairs of the structure: for a given (locally) admissible configuration, it will attempt to compute a satisfying (locally) healthy configuration. If it fails—having encountered a configuration that is not admissible, or a new burst—then the *rebuilding* procedure is called, which is designed to repair a larger interval. On a higher level of simulation, this corresponds to the implementation of the “cleaning” trajectory properties. The healing procedure runs in $O(\beta)$ steps, whereas rebuilding needs $O(Q^2)$ steps.

Peter remark 5. maybe even Q^3 ? ┘

The description of the procedures will look as if we assumed that there is no noise or disorder. The rules described here, however (as will be proved later), will clean an area locally under the appropriate conditions, and will also work under appropriately moderate noise.

The healing procedure does not even protect itself against any possible noise during it. The only protection is that any one call of the healing procedure will

change only a small part of the tape, essentially one cell: so a noise burst will have limited impact even if it happens during healing. One possible outcome of the healing procedure is *failure*: this happens when it encounters an unstitchable situation. Failed healing creates a *germ* of cells marked for rebuilding. The size of the germ is larger than $3\beta B$, to make sure that it could not have been created by a burst. Proceeding conservatively, the healing procedure carries out only one step of this plan, then it calls itself again. Even at failure, it will only add one cell at a time to the germ.

Where to put the germ? If no substantial domains are found then the germ is started in the middle of the healing area, and later steps (if still no substantial domains are found) add to it. If a substantial domain is found then the germ is started in the middle of one. This way, the next attempts will not try to eliminate it by stitching, and will add to it on failure. If healing succeeds then it ends with eliminating any possible started germ.

Recall the parameters Δ, E introduced in Definition 9.4. Suppose that *Heal* is called at some position z . Then it sets

$$Mode \leftarrow \text{Healing}, \text{Heal.Sweep} \leftarrow 1, \text{Heal.Addr} \leftarrow 0.$$

It starts by surveying an interval R of at least $2E$ cells around its starting point z . (It will go a little farther than E cells, in search of an allowed turning point, due to feathering. If such a point is not found within 3β steps, healing fails.) Whenever the head steps on a stem cell or creates a new cell, *Drift* is set to point to z (to make sure that the head does not get lost in a homogenous domain of stem cells).

Suppose that in the survey a pair of substantial domains separated by an ambiguous area is found, that is at least Δ removed from the complement of R . Choose the ambiguous area closest to the center (of healing). Then the first needed “stitching” operation (a single step) as defined in Lemma 9.6 is determined, and is performed. If the ambiguous area still remains, go to its cell closest to z and restart healing. In case that we started from a clean admissible configuration, this operation creates a new admissible configuration that is one step closer to being healthy.

If the required stitching operation is not found then the surveyed area is found inadmissible, and healing fails. Otherwise, if the head is at the front, healing is declared completed. Otherwise, the healing center is moved one step closer to the front, and healing is restarted.

The healing operation defined this way has the following property.

Lemma 9.7 *Assume that the head moves in a noise-free and clean space-time rectangle $[a, b) \times [u, v)$, with $b - a > 3EB$, $v - u > 2ET$, touching every cell of $[a, b)$ at least once, and never in rebuilding mode. Then at time v , the area $[a + 1.5EB, b - 1.5EB)$ is healthy.*

Proof. In healing, the head will not move away from an ambiguous area before eliminating it, creating a healthy interval I containing the two substantial intervals that originally bordered it. If the head leaves this interval in normal mode and the zigging does not start new healing then the healthy interval covered by zigging can be added to I to form an even larger healthy interval, and this continues until new healing starts. One of the substantial intervals of the new healing will be inside I , and when it is completed, I will be extended further. The head may later return into the healthy interval I , but it will then just continue the simulation without affecting health while staying inside. \square

9.3 Rebuilding

If healing fails, it calls the rebuilding procedure. This indicates that the colony structure is ruined in an interval of size larger than what can be handled by local healing. Just as with healing, we will be speaking here only about a situation with no mention of disorder—but the final analysis will take disorder into account.

Since rebuilding makes changes on an interval of the length of several colonies, it is important not only that it is invoked when it is needed, but that it is not invoked otherwise!

Rebuilding starts by extending the germ to the left, in a zigging way, expecting rebuild-marked cells on the backward zig. Since the zig is larger than a burst, if rebuilding started from just a burst then this will trigger healing, thus leaving the rebuilding mode.

The notion of front, of Definition 7.3 will be extended also to the rebuilding mode. Here is an outline, for rebuilding that started from a cell z in the middle of the germ. The goal is that

- the process does not destroy any healthy colony more than $3\beta B$ to the left of z or to the right of z .
- we end up with a new decodable area (see Definition 10.1) extending at least one colony to the left and one colony to the right of z .

The rebuilding operation uses its own addresses, but in a special way, to avoid being confused by the occasional deletion or insertion of cells. So it uses two

address tracks: $Rebuild.Addr_{-1}$ counts from the left end of the rebuilding area towards the head, and $Rebuild.Addr_1$ from the right end. Recall the stitching operation from Section 9.1.

Mark Starting from the germ, extend a rebuilding area over $3Q$ cells to the left and $3Q$ cells to the right from z . Mark the area using the tracks $Rebuild.Addr_j$ for $j = \pm 1$, and the track $Rebuild.Sweep$. False rebuilding started by a burst will be recognized since zigging must see an at least germ-sized marked area. In what follows, every step that changes the configuration must be accompanied by zigging, to check that the rebuilding is indeed going on.

Rebuilding cannot assume anything about the content of the area encountered: in particular, it may represent the traces of some other rebuilding. Since a single burst (which cannot be completely excluded) could pass control to this other rebuilding process, this possibility must be handled. If the front of the marking process encounters the front of a marking process in the opposite direction, then the marking process moving right gets precedence (see Remark 9.8 below), and so the area being marked in the left direction gets added to the right-directed one.

There is only one way that the rebuilding process can fail in a clean area: *turn-starvation*, as defined in Section 7.1.2.

Survey and Create More details of this stage will be given below. It looks for existing colonies (possibly needing minor repair) in the rebuilding area, and possibly creates some. As a result, we will have one colony called C_{left} on the left of z , one called C_{right} on the right of z , and possibly some colonies between them. Make all newly created colonies represent stem cells. Declare C_{left} the base colony, direct all the others with drifts and bridges towards it. (The creation of a bridge may result also in the creation of a new colony if the bridge becomes Q cells long.) The interval covering C_{left} and C_{right} will be called the *output interval* of rebuilding.

Mop Remove the rebuild marks, shrinking the rebuilding area onto the left end of C_{left} .

Remarks 9.8 1. Giving precedence to the right rebuilding has the drawback that one can design a initial configuration in which even in the absence of noise, level 1 structure will never arise even locally. Namely, we can fill the line with short intervals $\dots, J_{-1}, J_0, J_1, \dots$ each of which is the start (say of size $3Z$) of a right-directed marking process. Then the head, after moving left on J_0 , will be captured by the process on J_{-1} , then later captured by the

similar process on J_{-2} , and so on. But we will not need to consider such pathological configurations.

┘

Details of the Survey and Create stage

- s1) Search in the marked area on left of z for a colony C , or a set of cells that looks stitchable by up to 3 healings into a colony. The first substantial domain should be at least 3β cells to the left of z , to make sure that C is *manifestly* to the left of z . If the search and the stitching attempt are successful, mark C as C_{left} .

Repeat this search for a colony manifestly to the right of z : if found, call it C_{right} .

- s2) Suppose that only C_{left} is found, then create C_{right} . The other case is symmetrical.
- s3) Suppose that neither C_{left} nor C_{right} have been found. Then search the whole area, starting from the left, for a colony. If no such colony is found then create C_{left} and C_{right} .
- s4) Suppose that both C_{left} and C_{right} have been determined (found or created). Then search between them, from the left, for (stitchable) colonies, one-by-one.
- s5) Suppose that only a colony has been found that is not manifestly to the left or right of z ; then either the left end is manifestly to the left or the right end is manifestly to the right of z .

Suppose that the left end it is manifestly to the left of z , then call the colony C_1 . Then create C_{left} on the left of C_1 . Now search for another one on its right (it is not manifestly on the right). If not found, create C_{right} , manifestly on the right of z . If found call it C_2 , and create C_{right} on its right. The other case is symmetrical.

The steps above requiring the creation of colonies and bridges are destructive (the stitching ones are not). Therefore before a creation step, the whole survey preceding it is performed twice, marking the result in all rebuilding cells, on two different tracks. The required actions (erasing cells in order to build new ones in their place) are then only performed if the two survey results are identical—otherwise alarm is called (followed probably by new rebuilding).

Marked cells from some interrupted rebuilding may remain even after the mop-up operation. These may trigger new healing-rebuilding later.

10 Scale-up

10.1 Super-health

Definition 10.1 (Super-health) Let ξ be a clean configuration on an interval I . We say that ξ is *super-healthy* if both ends of I coincide with the end of a colony, and in each colony in I , whenever the head is not in the last sweep, the *Info* track contains a valid codeword as defined in Section 5.2. A tape configuration ξ is *super-healthy* if it can be extended to a super-healthy configuration. \lrcorner

As indicated in Section 6, when dealing with the behavior of machine M over some space-time rectangle, we will assume that the noise over this rectangle is $(\beta(B, T), \gamma(B^*, T^*))$ -sparse. With Definition 7.1 of T^* this means in simpler terms that at most one noise *burst* affecting an area of size at most βB can occur in any γ consecutive work periods. In the present section, histories will always be assumed to have this property.

Some histories lend themselves to be viewed as a healthy development that is disturbed only in some well-understood ways. Annotation adds to such histories the information pointing out these disturbances. The proof of the error-correcting behavior of machine M (essentially the proof of the Transition Function property of trajectories of Definition 6.10 for the simulated machine M^*) will take the form of annotation under sparse noise. An annotation, as per Definition 9.1, marks some ways in which the health of a tape configuration has been affected. Now we extend annotation in order to deal with damage not only to health but also to information.

Definition 10.2 (Super-annotation) A *super-annotated configuration* is a tuple

$$(\xi, \chi, \mathcal{I}, \mathcal{S}),$$

with the following meaning.

(ξ, χ, \mathcal{I}) is an annotated configuration.

\mathcal{S} is a set of disjoint intervals called *stains*. All islands are contained in stains.

We can change χ into a super-healthy configuration by changing it only in the stains. \lrcorner

Recall Definition 9.2 of admissibility.

Definition 10.3 (Super-admissibility) A super-annotation is *super-admissible* on an interval K if it is admissible on K , further consider any interval $J \subseteq K$ of size $\leq QB$.

a3) At most 3 stains intersect J .

- a4) If J contains k stains, surrounded by some healthy area (we already know $k \leq 3$), while the head is at a distance $> 2B$ within the clean area, then the total size of these stains is at most $k \cdot c_{\text{stain}}B$, where

$$c_{\text{stain}} = (f + 2)\beta. \quad (10.1)$$

- a5) If an outer colony in K intersects two stains then it simulates a cell state with $Turned = 1$.

┘

Requirement (a5) is essentially motivated by the consideration described in Example 9.3. So a second stain can remain at the end of a work period in which the simulated machine makes a turn, setting $Turned \leftarrow 1$ in the simulated cell. Requirement (a5) says that in an annotated configuration, this is the only way for two stains to remain in an outer colony.

A configuration may allow several possible super-annotations; however, the valid codewords (referred to in the definition of super-health) recoverable from it do not depend on the choice of the annotation.

The following is an immediate consequence of the definition of rebuilding: in a clean area and the absence of noise, rebuilding will succeed.

Lemma 10.4 *Suppose that a rebuilding procedure starts on the boundary of a super-admissible subinterval $J \subset I$ of a clean interval I , and runs noiselessly to the finish. Denoting by K the output interval of rebuilding, then $J \cup K$ will be super-admissible.*

10.2 Annotated history

The following definitions extend the notion of annotation to histories, in order to discuss the effects of moderate noise and their repair.

Definition 10.5 (Distress and relief, safety) Consider a sequence of annotated configurations over a certain time interval. If the head is free (see Definition 9.1), then the time (and the configuration) will be called *distress-free*. A space-time point that is not distress-free and is preceded by a distress-free time will be called a *distress event*. This can be of two kinds: the head steps onto an island, or a burst occurs (creating an island and leaving the head in it).

Consider a time interval K starting with a distress event and ending with a distress-free configuration. Let J be the interval of tape where the head passed during K , then we will call $J \times K$ a *relief event*, if the following holds.

- a) Any new islands occurring in $J \times K$ are due to some new burst.

b) The island that started the distress event disappears by the end of K . ┘

In a relief event, it is possible to leave behind some islands other than the one initiating the distress.

We will consider the annotation of histories over a limited space-time region, but will not refer to this region explicitly every time.

Definition 10.6 (Annotated history) An *annotated history* of a generalized Turing machine

$$M = (\Sigma, \tau, Adj, B, T, \pi)$$

is a sequence of super-annotated configurations such that the sequence of underlying configurations is a trajectory, and every distress event is followed by a relief event $J \times K$ with $|J| = O(\beta B)$ and $|K| = O(\beta^2 T)$. ┘

In what follows we will show that for any trajectory $(\eta, Noise)$ of a generalized Turing machine M on any space-time rectangle on which the noise is $(\beta(B, T), \gamma(B^*, T^*))$ -sparse, if at the beginning the configuration was super-healthy then the history can be annotated. In the rest of the section we always rely on the assumption of this sparsity property of the noise.

The main part of the proof is about obtaining relief after a distress event. Unlike in [1], now islands may have their cell structure damaged: may contain disorder. However, since η is a trajectory, as we will see the islands will be cleaned out. So, relief maybe said to happen in two stages: cleaning, and correcting the structure. However, this division is only for the purposes of proof: the machine has no “disorder-detector”, and the proof just relies on the cleanness-extending properties of a trajectory introduced in Definition 6.10.

10.3 The simulation codes

Let us now define formally the codes φ_{*k}, Φ_k^* needed for the simulation of history $(\eta^{k+1}, Noise^{(k+1)})$ by history $(\eta^k, Noise^{(k)})$. Omitting the index k we will write φ_*, Φ^* . To compute the configuration encoding φ_* we proceed first as done in Section 6.3, using the code ψ_* there, and then initialize the sweep and drift and address fields appropriately.

The value $Noise^{(k+1)}$ is obtained by a residue operation just as in Definition 5.2 of sparsity. It remains to define $\eta^* = \eta^{(k+1)}$ when $\eta = \eta^k$. Parts of the history that are locally super-annotated will be called *clean*. In the clean

part, if no colony has its starting point at x at time t , set $\eta^*(x, t) = \text{Vac}$. Otherwise $\eta^*(x, t)$ will be decoded from the *Info* track of this colony, at the beginning of its work period containing time t . More precisely:

Definition 10.7 (Scale-up) Let (η, Noise) be a history of M , where $\text{Noise} = \text{Noise}^{(k)}$ as in Definition 5.2. We define $(\eta^*, \text{Noise}^*) = \Phi^*(\eta, \text{Noise})$ as follows. Let $\text{Noise}^* = \text{Noise}^{(k+1)}$. Consider position x at time t , and $J = (t - T^*, t]$. If x is not contained in some interval I such that η is super-admissible over $I \times J$ then $\eta^*(x, t) = \text{Bad}^*$. Assume now that it is contained, and let $\chi(\cdot, u)$ be some super-healthy history satisfying η over $I \times J$. If x is not the start of a colony in χ then let $\eta^*(x, t) = \text{Vac}$; assume now that it is. Then let $t' \in J$ be the starting time in χ of the work period of C containing t , and let $\eta^*(x, t)$ be the value decoded from $\eta(C, t')$. In more detail, as said at the end of Section 5.2, we apply the decoding ψ^* to the interior of the colony to obtain $\eta(x, t)$. \dashv

This definition decodes super-admissible intervals and histories for η into super-healthy intervals and histories of η^* , while preserving the property (C1) of cleanness in Definition 6.3.

11 Isolated bursts

Here, we will prove that the healing procedure indeed deals with isolated bursts. Our goal is to show that it provides relief, as required in an admissible annotated trajectory. For the elimination of disorder created by bursts we will rely on the Escape, Spill Bound and the Attack Cleaning properties of a trajectory in Definition 6.10.

Isolated bursts don't create disorder larger than 3β . The head escapes a disorder interval I via the Escape property; while it is inside, the spreading of this interval is limited by the Spill Bound property. Every subsequent time when the head enters and exits I this gets decreased via the Attack Cleaning property, so it disappears after $O(\beta)$ such interactions—see Lemma 11.2 below.

In a clean configuration, whenever healing started with an alarm, the procedure will be brought to its conclusion as long as no new burst occurs. The trajectory properties, however, do not allow any conclusion about the state of the cell to which the head emerges from disorder. This complicates the reasoning, and may require several restarts of the healing procedure. By design, the healing procedure can change the *Core* track only in one cell. The following lemma limits even this kind of possible damage: it says that the head with the

wrong information may increase some existing island, but will not create any new one.

Lemma 11.1 *In the absence of noise, no new island will arise.*

Proof. The islands are defined only by the *Core* track. In normal mode, this track changes only at the front. If this is not the real front, then we are already in or next to an island.

The healing procedure's change of *Core* is part of a stitching operation. Looking at the different cases of the proof of Lemma 9.6, we see that inside a healthy area, healing can only change the *Core* track in two ways.

The first way is case 1, when the *Sweep* values were changed in a domain not belonging to the interior. Such an operation can be applied to any healthy configuration without affecting health.

The second case is when the front is moved left or right. This does not affect health either. \square

The following lemma is central to the analysis under the condition that bursts are isolated.

Lemma 11.2 (Healing) *In the absence of noise in M^* , the history can be super-annotated. Also, the decoded history (η^*, Noise^*) satisfies the Transition Function property of trajectories (Definition 6.10).*

Proof. The proof of super-annotation is by induction on time, extending the super-annotation into the future. The extension is straightforward as long as no distress event is encountered. The Transition Function property is observed, as no obstacle arises to the simulation. Stains do not cause problems: the computation stage of the simulation cycle eliminates them using the error-correcting code.

Consider now the occurrence of a distress event. This can be due to either a burst or the encounter with an island. In the latter case, before the relief a burst can still hit. In the analysis below, then we will just cut our losses and restart, knowing that burst cannot hit again before relief.

At the time of the distress event, let us draw an interval I of size QB centered around the head. The head will not leave it before relief, so we will consider the changes of the configuration inside it. Let S_1, \dots, S_m be the list of substantial domains of I , and A_1, \dots, A_{m-1} the ambiguous domains between them. Note that $m = O(1)$. Let R be the set of those i for which $S_i \cup A_i \cup S_{i+1}$ is clean. For $i \in R$ let n_i be the number of stitching steps by the algorithm of Lemma 9.6

needed to stitch them together. Let $N = \sum_{i \in R} n_i$. Whenever the head enters disorder, we will mark the edge where it entered as *attacked*.

We introduce a few variables for the proof.

- D = be the total size of disorder, divided by B .
- E = the number of un-attacked edges of disorder where the head can exit without entering (because it is on the side of the disorder). This number never increases.
- F = the number of un-attacked edges where the head cannot exit without entering (because it is away from the disorder).
- P = be the number of steps that the front moves forward (including the ones after possibly changing the meaning of forward at a regular turn).

The following kinds of event may occur:

- h1) With a stitching done, N decreases by 1 while possibly moving the front backward.
- h2) The head enters disorder, decreasing F .
- h3) The head leaves disorder on a non-attacked edge, decreasing E .
- h4) The head leaves disorder on an attacked edge, decreases the disorder by at least B and increases F by 1.
- h5) A set A_i becomes clean: then i gets added to R ; this can happen only m times.
- h6) N increases in the healing mode by 1. This can only happen after the head exited disorder in healing mode (with the wrong information). At the exit, either E or D had to decrease.
- h7) The head approaches the front: if it does not reach there then it gets closer by $\Omega(\beta)B$.
- h8) The front moves forward in normal mode (possibly after making a regular turn, and thus changing the forward direction). This may also increase N by 1 (and is the only way to do so), say by decreasing a domain S_i . But it is followed by zigging. The zigging may hit disorder, leading to case (h2), or find something wrong—and trigger new healing, which leads to some of the other events. Otherwise the island around the front will be deleted, and the head becomes distress-free.

The above possibilities suggest a potential function

$$N + c_D D + c_E E + c_F F + c_m m - c_P P$$

where c_D, \dots are appropriate positive constants. Let us look at what each possibility does to the potential.

Cases (h1-h3) decrease the potential if $c_P < 1$. Case (h4) decreases the potential if $c_D > c_F$. There are only $O(1)$ cases of type (h5), increasing N by a total of $O(\beta)$.

In case (h6), if c_D and c_E are large enough, the increase in N can be charged to a decrease in D or E .

Case (h7) will not change the potential but there are at most $O(1)$ consecutive such cases.

Consider case (h8). If the zigging hits new disorder and $c_F > 1$ then the potential decreases. If it triggers new healing then this will lead to one of the other cases, compensating for the increase of N if the constants (other than c_P) are large.

These considerations show that relief indeed follows distress in time $O(\beta^2 T)$. At that point a normal-mode step moving the front has been made, followed by zigging. Therefore the island causing the distress must have been eliminated, allowing the simulation to continue. The stain caused by the island remains, but as discussed after Definition 10.3, the simulation guarantees that the size and number of stains remains bounded as required by that definition. The simulation also continues to satisfy the Transition Function property of trajectories. \square

12 Cleaning

This section will scale up the Spill Bound, Escape, Attack Cleaning and Pass Cleaning properties of trajectories.

12.1 Cleaning the current level

We need some lemmas in preparation.

The proof of the following lemma does not use any properties of the simulation program.

Lemma 12.1 *Let P be a space-time path of the head that has no bursts, and makes at least $\pi + \theta$ pairs of passes over an interval J_1 of size $\theta B/2$. Then there is a time during P when $\text{Int}(J_1, c_{\text{pass}} B)$ becomes clean.*

Note that the assumption of the lemma do not bound the number of θB -intrusions.

Proof. Assume that J_1 is an interval with the required properties that does not become clean at any time during P ; we will derive a contradiction. Let t_1 be the end time of the first π pairs of passes of P over J_1 .

Since J_1 did not become clean, the number of θB -intrusions into it before t_1 is more than π^2 . The ones that come from left will be called *left intrusions*. Without loss of generality we can assume that the number of left intrusions is

$$K_2 > \pi^2/2,$$

coming in $\pi^2/4$ pairs. Let t_2 be the time before the last θ pairs of left intrusions. Let J_2 be the interval of length $|J_1|$ adjacent to J_1 on the left. Since the intrusions don't reach the right end of I but have length $\geq \theta B$, they extend to a distance $\theta B/2$ to the left of I , so each of them passes J_2 , giving at least $\pi^2/4 - \theta$ pairs of passes over it. If J_2 becomes clean before time t_2 then the next θ passes over J_1 clean interval J_1 via the Attack property, contrary to our assumption.

Let us iterate this reasoning for $i = 2, \dots$. Assume we have K_i pairs of passes over J_i before time t_i . Let t_{i+1} be the time $t_{i+1} < t_i$ before the last θ pairs of them. We can divide the ones before t_{i+1} into groups of $\pi + \theta$ consecutive pairs of passes, so we have now at least

$$K_i - \theta$$

pairs of passes over J_i before t_{i+1} divided into consecutive groups of size $\pi + \theta$. Each group fails to clean J_i , so the number of its θB -intrusions is more than π^2 , giving a total of more than

$$\pi^2(K_i - \theta)/(\pi + \theta) > 4\pi(K_i - \theta)/5$$

if π is large. Without loss of generality, at least half of these intrusions is on the left of J_i . Let J_{i+1} be the interval of size $|J_i|$ adjacent to J_i on the left, so just as in the reasoning about J_2 , we conclude that it is passed over in at least

$$K_{i+1} = \pi(K_i - \theta)/5 = K_i\pi/5 - \theta\pi/5 \quad (12.1)$$

pairs. Hence for $i \geq 3$

$$K_i = (\pi/5)^{i-1}\pi - \theta((\pi/5) + (\pi/5)^2 + \dots + (\pi/5)^{i-2}) \geq (\pi/5)^i$$

if π is large. Since P is finite, sooner or later this process must terminate, leading to a contradiction. \square

12.2 Escape

This section scales up the Escape property of trajectories. We will frequently measure the length of a path not by the time it takes but its number of segments, in the sense of Definition 6.12.

Definition 12.2 In the present section, a *segment* we will mean a θB -segment. ┘

Note that one work period of simulation or rebuilding happens within an area of size $\leq \theta QB$. If a path is burst-free then the Escape property allows to estimate the time length of the path by its number of segments:

Lemma 12.3 *A burst-free path with s segments takes time $\leq c_{\text{esc}}\theta^2(s+1)T$.*

Proof. Given that $\theta < \beta$, The Escape property says that the path leaves every segment of size $\leq \theta B$ in time $c_{\text{esc}}\theta^2T$. □

In what follows we will estimate the length of paths mainly by their number of segments.

Definition 12.4 For an interval I , the interval $\text{Int}(I, c_{\text{spill}}B)$ be called its *spill-interior*. A maximal clean interval of size $\geq c_{\text{attack}}B$ will be called a *clean hole*. Let $\mathcal{K}(t)$ denote set of all clean holes at time t , let $K(t) = \bigcup \mathcal{K}(t)$. Let $\mathcal{K}^*(t)$ be the set of super-healthy clean holes, and $K^*(t) \subseteq K(t)$ be their union. ┘

Consider a burst-free space-time rectangle $I \times J$. The Spill Bound property implies that no disorder can appear during it in the spill-interior of a clean hole. If I_1 is a clean hole at time t and I_2, I_3 are clean holes at time $t+1$ intersecting the spill-interior of I_1 then the smallest interval containing $I_2 \cup I_3$ is contained in a clean hole at time $t+1$. So the elements of the set $\mathcal{K}(t)$ of clean holes of size $\geq c_{\text{attack}}B$ will not break up: they can grow, shrink but only slightly, bounded by the Spill condition (and thus will not disappear) or merge; new elements can appear as well. They may intersect without merging, but only to an extent less than $c_{\text{spill}}B$. We can similarly follow a single clean hole $I(t)$.

Just as a clean hole cannot break up in the absence of bursts, a super-healthy interval cannot break up either. So we can follow the evolution of super-healthy clean holes in $\mathcal{K}^*(t)$ similarly. Recall that a super-healthy interval in history η is clean in the decoded history η^* , so the simulation of η^* will proceed in it. In what follows we will estimate how much the area of the sets $K(t)$ and $K^*(t)$ will grow with the cumulative time spent by the head in $K(t)$. The following lemmas consider the activity of the head while it is inside a clean hole $I(t)$.

If the head stays in long enough then it will create colonies if they were not there.

Lemma 12.5 *Let $I(t)$ be a clean hole having size $\leq nB$ with $n \leq \lambda Q$, $\lambda \leq 3\beta$ at the beginning of a burst-free path P , and $k_I(t) = |K^*(t) \cap I(t)|$. If $n < Q/2$ then the head will leave $I(t)$ within cZn segments for an appropriate c . Otherwise, every λU segments of its stay increases $k_I(t)$ (that is the size of area “organized up to the next level”) by at least QB .*

Proof. We will see that while the head does not leave $I(t)$, every certain number of segments results in some kind of progress.

1. Suppose that we are at some time when the rebuild procedure had started, from a base of rebuild-marked cells large enough not to result in alarm (renewed call for healing) after the first zigging. Then the rebuilding will be completed.
2. Suppose that we are at some time when the mode is normal. Then within $2Z$ steps of the simulation (and hence at most $2Z$ segments) either a healing call happens, or a step of progress will be made in the ordinary work of simulation.
3. Suppose that at some time healing will be called. Then according to the proof of Lemma 11.2, we either arrive at the above case 1 (failed healing) in $O(\beta^2 < Z)$ steps or healing finishes in normal mode with at least one step made either to move closer to the front or to make a move in the ordinary work of simulation.

Indeed, if healing succeeds it leaves a healthy area. Lemmas 8.7 and 9.7 imply that the overlapping successful healing areas can be combined, so healing will not be repeated over the same area, and thus can slow down progress only by a factor $O(\beta^2)$ (negligible compared to the Z times slowdown by zigging), and even this in at most one sweep of simulation over any area.

4. Suppose that $n > Q/2$ and the head stays for more segments than λU . Without rebuilding, the continued healings add some colonies to the healthy area; any such colony will be turned super-healthy by the first complete work-period of the simulation. If the head is already in a super-healthy area then it will perform the simulation of M^* . The simulated machine M^* has the same program as M , so it will also make switchbacks of at least $E > 3\beta$ steps (for healing) or even bigger ones for zigging. This extends the super-healthy area unless interrupted by rebuilding.

5. If any rebuilding starts then it will either finish and extend $K^*(t)$ by a colony, or will be interrupted by a turn-starvation, as defined in Section 7.1.2. Suppose that turn-starvation happens at an attempted left turn: then the marked healing area is already of a size $\geq c_{\text{turn-limit-2}}Q$. Therefore a next turn starvation cannot be due to an attempted right turn on the left, since the marking process does not proceed so far to the left. Repeated turn-starvations could only happen to the right, and the head would leave on the right before the assumed λU segments of stay. Since this does not happen, rebuilding must succeed before this.

□

Peter remark 6. Elaborate!

┘

Now we consider cases where the head exits and enters a clean hole repeatedly.

Definition 12.6 Let I be a clean interval. We will say that the right end of I is *advancing* at time t if the rightmost colony of $K^*(t) \cap I$ is in the start of the process of transfer to the right. Advancing is defined similarly for the left end. ┘

Lemma 12.7 Consider the conditions and notation of Lemma 12.5.

- a) If the head stays in $I(t)$ longer than cZn segments and $k_I^*(t)$ did not increase by at least QB , then the end on which it leaves becomes advancing.
- b) If the head enters $I(t)$ on an advancing end and stays longer than cZn segments then $k_I^*(t)$ increases by at least QB .

Proof. To (a): If the head did not continue the simulation on $K^*(t)$ then it would have spent its time in healing-rebuilding. In this case the time is sufficient to create a new colony, increasing $k_I^*(t)$ by at least QB . Since this did not happen, it must have continued a simulation. But then it can leave $K^*(t)$ only via a started transfer operation.

To (b): Suppose that the head enters on an advancing right end. The head cannot get to the left of the colony C that started the right transfer operation without completing it, and entering C from the right from a new colony (either created by the transfer or by a new rebuilding operation). And it spends so much time on the right of C that at least the rebuilding operation must succeed. □

The next lemma considers a situation similar to Lemma 12.5, but allowing multiple entries and exits.

Lemma 12.8 (Expand hole) *Let $I(t)$ be a clean hole having size nB where $n = \lambda Q$ with $1/2 \leq \lambda \leq 3\beta$ at the beginning of a burst-free time interval J . There is a constant c with the following property. If the head spends cumulatively $c\lambda U$ segments in $I(t)$ then at the end of J the hole grows in size by at least $QB/2 - c_{\text{spill}}B > QB/3$.*

Proof. If the head spends $\leq cZn$ segments in $I(t)$, with the constant c from Lemma 12.7, then we say that its stay was *short*. If it spends $\geq \lambda U$ segments in $I(t)$, then we will say that the stay was *long*. Otherwise we will say that the stay was *medium*.

1. After a short stay (except possibly the first exit on the left and the first exit on the right), $I(t)$ expands by at least $B/2$ via the Attack property.
2. After a long stay (even without leaving), by Lemma 12.5, the value of $k_I(t) = |K^*(t) \cap I(t)|$ increases by at least QB .
3. Lemma 12.7 implies that if a medium stay did not increase $k_I(t)$ by at least QB then it creates an advancing end, and the next medium stay will produce such an increase.

Adding up these possibilities will complete the proof. \square

The following lemma is the scale-up of the Escape condition.

Lemma 12.9 (Escape) *There is a constant c such that in the absence of Noise*, the head will leave any interval G of size nB with $n \leq \lambda Q$, $1 \leq \lambda \leq 3\beta$, within cU segments.*

Proof. If J is the time interval considered, then we can ignore the at most one burst happening during J by considering the half of J where it does not occur.

At any time t , we will partition the interval G into subintervals as follows. First, let $\mathcal{K}(t)$ be the set of clean holes of size $\geq c_{\text{attack}}B$ inside G , and $K(t)$ their union. Each interval between two such clean holes that is larger than $\theta B/2$ will be subdivided into a sequence of subintervals; the last one has size $\leq \theta B/2$, the other ones have size $\theta B/2$. Let us call all members of this partition of G *blocks*; the clean holes among them will be called *clean*, the other ones are called *short*.

As time passes new clean blocks may arise, the existing ones may increase, and neighboring ones may merge. (They may also decrease due to spill of disorder but this decrease is temporary until the next increase due to attack.)

We will consider a sequence of times $t_0 < t_1 < \dots$ defined as follows. t_0 is our starting time. If t_i is defined and then the head is in some block S , then t_{i+1} is the first time coming at which

- if S is clean then the head is either not in S or the path entered a new segment.
- if S is short then the head is not in either S or in any short block adjacent to it.

By definition, between times t_i and t_{i+1} one of the following events occurs. We will show the contribution to the growth of $|K(t)|$ made by them. Clearly $K(t)$ cannot grow larger than $|G|$.

- e1) A boundary is crossed between clean and short blocks. A pair of passes over the same boundary results in an increase of $K(t)$ by B . The number of unpaired passes over boundaries is at most the upper bound of the number of such boundaries, that is e_2Q where

$$e_2 = \lambda/c_{\text{attack}}.$$

So if s segments are spent on events of this type then such events, and then the growth of $K(t)$ is at least

$$\frac{s - e_2Q}{2} \cdot \frac{B}{2}.$$

Assuming $s \geq 2e_2Q$ this is at least $sB/8$. This becomes larger than $|G| = \lambda QB$ if $s > 8\lambda Q$ so the number of segments spent on these events is less than this.

- e2) Two clean blocks merge. The number of these events is limited by e_2Q .
- e3) The number of segments spent in $K(t)$ increases by one. Let c_{expand} be the constant c of Lemma 12.8. By that lemma, if the head spends time $c_{\text{expand}}\lambda U$ segments on these events then $K(t)$ increases by at least $QB/3$. So it cannot spend more than $3c_{\text{expand}}\lambda^2U$ without escaping G .
- e4) The number of passes over a short block S increases by 1. If $\pi + 2$ pairs of passes happen over a short block S then Lemma 12.1 becomes applicable and $\text{Int}(S, c_{\text{pass}}B)$, of size $\geq \theta B/6$, becomes clean, increasing $|K(t)|$ by $\theta B/6$. If n is the number of such events then the increase is at least

$$\frac{n\theta B}{2\pi + 4}.$$

This becomes larger than λQB if $n > (2\pi + 4)\lambda Q/\theta$.

□

12.3 Pass cleaning

This section will scale up the Pass Cleaning property to machine M^* . We will consider a path P with no bursts of η^* , as it makes passes over an the interval I of size θQB . First a strengthening of Lemma 12.1: we allow bursts of η .

Lemma 12.10 *Let $r = 4$, and suppose that a path P has no bursts of η^* , it makes $\pi + \theta$ pairs of passes over the interval I of size θQB , with the number of θQB -intrusions at most $(\pi + 2\theta)^2$. Let $J_1 \subset \text{Int}(I, c_{\text{pass}}QB/2)$ be an interval of size θB at a distance $\geq r\theta B/2$ from any burst of η . Then there is a time during P when $\text{Int}(J_1, c_{\text{pass}}B)$ becomes clean.*

The proof still uses no properties of the simulation program.

Proof. 1. There are at most $2\pi + (\pi + 2\theta)^2 < 2\pi^2$ bursts of P in I , and at most $2c\pi^2U$ segments, where the constant c is the one from Lemma 12.9.

Proof. The Escape Lemma 12.9 implies that every pass of P over I leaves using at most cU segments, in time T^* , hence has at most one burst. Each θQB -intrusion lasts also at most time T^* . So the total number of bursts is at most the sum of the number of passes and intrusions. A bound on the total number of segments is obtained multiplying this number by cU .

2. Since the interval J_1 is at a distance $\geq r\theta B/2$ from any burst, we can repeat the reasoning of the proof of Lemma 12.1, for $i \leq r$. We again define a sequence of adjacent intervals J_i such that the path P makes $K_i \geq (\pi/5)^i$ pairs of passes over J_i
3. We will handle the case $i = r$ differently, using the observation that the total number of bursts is at most $s = 2\pi^2$.

Each burst may affect at most one pair of passes, so there remain at least $K_i - s$ burst-free ones of them, separated into at most $s + 1$ burst-free groups. At least one of these groups has size at least

$$K_{r+1} = \frac{K_r - s}{s + 1} > \frac{K_r - 2\pi^2}{2\pi^2 + 1} \geq \frac{\pi^{r-2}}{3 \cdot 5^r}.$$

if π is large. Let t_{r+1} be the endtime of the last pass in this group. With $r = 4$ this gives $K_5 \geq \pi^2/3 \cdot 5^4 > \pi$ if π is large.

4. From here on we can again use (12.1), since no burst appears at all during the time of this group of passes, and we get $K_i \geq (\pi/5)^{i-4}$. The number of θB -intrusions into J_i grows correspondingly. Since $J_1 \subseteq \text{Int}(I, c_{\text{pass}}QB/2)$, before J_i gets outside I we have $i > c'Q$ for some constant c' , giving a lower bound exponential in Q on the total number of segments.

But by part 1 above, the total number of segments in P is bounded by

$$2c\pi^2 U \leq 2c_{\text{sim}} c\beta\pi^2 Q^2$$

where we used the definition of U in Definition 7.1. Since this is only quadratic in Q , the lower and upper bounds collide, delivering the desired contradiction. □

Under the conditions of Lemma 12.10, the interiors of burst-free subintervals of I become clean. By the Spill Bound property, they essentially also remain clean as long as they remain burst-free. The following two lemmas will help expanding them.

Recall that $E \ll Z$ is the maximum number of cells in a healing area.

Definition 12.11 An interval J not containing the head is called *right-directed* if the head is to the right of J , and its cells, maybe with the exception of an island and an area of size EB on the left end and one of size kEB with $k < Z/2$ on the right end, point towards a front at the right end of J (in normal operation or rebuilding), with the corresponding frontier zone. We will say that J has *damage size* k . ┘

Lemma 12.12 *If the head passes through a clean interval J of size $> 4ZB$ noiselessly from left to right then J becomes right-directed with no damage.*

Proof. The head is acting in one of three modes: normal, healing or rebuilding. Assume that the starting mode is normal. If it remains normal then J naturally becomes directed by the time the head leaves. If healing mode gets triggered then as long as healing succeeds it extends a healthy area. It moves towards right only if the front is on the right.

If it does not succeed then rebuilding gets triggered, and since it does not touch the left end anymore, it either succeeds or exits on the right while trying to extend towards the right. The same considerations work if the starting mode is healing, except possibly on the part of the left end where the first healing took place, whose purview was not completely contained in J .

The analysis is similar if the starting mode is rebuilding. □

Recall the definition of the parameter F in (7.4).

Lemma 12.13 *Suppose that a basic interval J is right-directed, the head enters it and leaves it, with at most one burst happening during this.*

a) *If the head leaves on the right then J will stay right-directed.*

b) *If the head leaves on the left then the right end is a legitimate turning point of the simulation or the rebuilding (that is J has almost F cells with $\text{Turned} = 2$ on its right).*

Proof. To (a): Both in normal operation and rebuilding, the head moves the front with itself and returns to it from every zig, except when it is captured at an end of J (by a burst or disorder). If an island is encountered or a burst occurs then it will be healed, except when the healing interval (whose maximum size is EB) intersects the boundary of J (where the disorder may capture the head).

To (b): the front can turn back only at the legitimate turning points; otherwise the simulation will move it forward, and the forward zigging prevents a burst or disorder from turning it back prematurely. \square

Let

$$\pi^* = \pi + 2\theta. \quad (12.2)$$

Lemma 12.14 (Pass cleaning) *Suppose that a path P has no bursts of η^* , it makes π^* pairs of passes over the interval I , and the number of θQB -intrusions is $\leq (\pi^*)^2$ for I and P . Then by end of the π^* th pass the interior $\text{Int}(I, c_{\text{pass}}QB)$ becomes clean for η^* .*

Proof. Just as in the proof of Lemma 12.10, there are at most

$$\kappa = 2\pi^2 \quad (12.3)$$

bursts of P in I . That lemma shows that at some time t_0 after $\pi + \theta$ passes, the areas of I at least a constant distance away from the bursts become clean, so the interval

$$I' = \text{Int}(I, c_{\text{pass}}QB/2)$$

becomes clean for η , except for $O(\pi^2)$ islands. So the interval I contains clean subintervals of size $\geq 4ZB$ separated from each other by areas of size $O(\pi^2 ZB)$. Let us call these clean subintervals *basic holes*.

1. Two pairs of passes make $I' \setminus \text{Int}(I', O(\kappa E))$ clean.

Proof. Lemmas 12.12 and 12.13 show that

- Passing a basic hole left to right makes it right-directed.
- The limited number of intrusions keep it right-directed with the damage staying bounded by $\kappa E \ll Z$.

- Passing a right-directed basic hole right to left can happen only if the right end is a legitimate turning point (of the simulation or rebuilding), even if a new rebuilding was triggered at the right end.
- Feathering does not allow the following right pass over the same basic hole to turn back before going a distance at least FB to the right. As $F \gg \kappa Z$, which is the upper bound on the separation between the basic holes, the remaining disorder between them must be wiped out before the head can move left over them in the second pair of passes.

2. The following pass makes $\text{Int}(I', O(\kappa E))$ clean for η^* .

Proof. Before the last pass from right to left, there can be several intrusions. Each such intrusion will turn back only at a legitimate turning point. It may leave an island if a burst happens on its left end. If a next intrusion passes over this island then due to feathering it must pass at a significant distance, and clean away the island. If it turns back before passing it then it must turn back at a distance of FB , so the islands left this way are separated by FB . When finally a left pass comes it can only pass over the whole interval I while cleaning all these islands and finding or rebuilding whole colonies.

□

Remark 12.15 Left and right are not symmetric in the above proof. Consider a right-directed basic hole with a rebuilding area being extended on the right. The head may be captured and returned, extending some other rebuilding area towards the left. Without giving preference to the right direction, this switching between left and right-directed rebuilding could happen an unbounded number of times. Using the asymmetry of the definition of rebuilding in Section 9.3, the right-directed extension of the basic hole will not be overridden by a left-directed rebuilding. ┘

12.4 Spill bound

Let us prove the scaled-up version of the spill bound property.

Lemma 12.16 *Suppose that an interval I of size $> 2c_{\text{spill}}QB$ is clean for η^* . and let P be a path that has no bursts of η^* . Then $\text{Int}(I, c_{\text{spill}}QB)$ stays clean for η^* .*

Proof. The decrease of I by $c_{\text{spill}}QB$ at any of the edges is possible due to some rebuilding process starting there. Due to the preference to the right expansion of rebuilding, there can be only one rebuilding area on the right edge of I , but the left edge may see a sequence of overlapping rebuilding areas. Since this is the

more complex possibility, let us just concentrate on this one, so all the considered exits and entries of the path into I will be on the left.

1. The lack of bursts of η^* and the escape property implies that each incursion of the path into I contains at most one burst. For simplicity, we will always assume that there is a burst, referring to it as *the* burst.

When the path first enters into I , which is clean for η^* and therefore healthy for η , then if it can work in normal mode, it will continue the simulation. The burst will be healed unless it is either at a legitimate turning point or is closer than EB to the edge of I . In what follows the clean part I' of I may become smaller and may contain islands but only at legitimate turning points. All colonies except the leftmost and rightmost ones, may contain only one island. Indeed, if the head passes those in the normal course of simulation then only the last pass can leave an island.

The path will enter rebuilding mode only if it is within EB of the edge of I , and a healing did not succeed (or even start). In this case, it will extend a rebuilding area J_1 starting from near the left edge. If the rebuilding succeeds then it will extend the healthy area of I on the left, and we are done.

But the path may leave it on the left edge before finishing. It may in this case leave a single island if the burst happened at some legitimate turning point, or just as above, near the left edge of I . When it enters next time, it may not continue the work of rebuilding of J_1 but may also start a new rebuilding area J_2 on the left edge of J_1 that may overwrite part of J_1 due to the priority of right-extending rebuilding. It will heal any island from the previous incursion that it encounters, since the place of this island cannot be a new turning point now. So now the picture is a disjoint union $I = D \cup J_2 \cup J_1 \cup I'$, where D is the possible union of islands formed by the bursts at the left edge, and I' is clean.

When the head enters next time, it can only work on J_2 (possibly extending it over J_1) or start a new rebuilding area J_3 on the left edge of J_2 .

Generalizing this picture we say that I is *defended on the left* if

- d1) It is a disjoint union $I = D \cup J \cup I'$ where I' is clean for η^* , further

$$J = J_1 \cup J_2 \cdots \cup J_k$$

where each J_i is a rebuilding area.

- d2) The endpoints of J_i are at legitimate turning points and therefore at a distance of at least FB from each other.
- d3) There is at most one island in each J_i , also at legitimate turning points and therefore also at a distance of at least FB from all other islands.

d4) There is also at most one island introduced by P in I' .

The above reasoning shows that the interval I will always stay defended on the left. Moreover, if J is non-empty then it can only become empty again when all the rebuilding on the left succeeds. As we will see then we will be done, since below we will bound $|D|$. If J stays non-empty then also, since $|J \cap I|$ is at most as large as the largest possible rebuilding area, we only need to bound $|D|$ to be done. For this, we will bound the total number of bursts created by P in I by 2π .

The lack of bursts of η^* implies that between consecutive bursts the path must run long stretches in time. Let us show that it must spend these stretches (except for a couple) on the left.

2. If J stays empty then P cannot enter more than $3QB$ deep into I without extending clean area of I on the left.

Proof. Let C be the leftmost colony of I not involved into a rightward transfer operation (so it is the leftmost or second leftmost colony of I). If P passes C from right to left then next time it can enter C only from another colony on its left.

Let I_0 be the interval of size θQB on the left of I . From the above it follows that every time (but maybe two) when P leaves a burst, it must make a pass over I_0 . Now an argument similar to the proof of Lemma 12.1, scaled up to η^* , shows that at some time during these passes, $I'_0 = \text{Int}(I_0, c_{\text{pass}}BQ)$ becomes clean for η^* . But then the subsequent passes from I'_0 into I will clean the area between them and unite them. \square

12.5 Attack cleaning

This section will scale up the Attack Cleaning property of trajectories (Definition 6.10) to machine M^* . First we prove a weaker version, requiring the space-time rectangle of interest to be free not only of Noise^* but also of Noise —that is burst-free.

Lemma 12.17 *Consider a trajectory η in a noise-free space-time rectangle. Here, the decoded history η^* satisfies the Attack Cleaning property.*

Proof. The property says the following for the present case. For current colony-pair x, x' (where $x' < x + 2QB$), suppose that the interval $[x - c_{\text{attack}}QB, x' + QB]$ is clean for M^* . Suppose further that the transition function, applied to $\eta^*(x, t)$, directs the head right. Then by the time the head comes back to $x - c_{\text{attack}}QB$, the right end of the interval clean in M^* containing x advances to the right by at least QB .

The computation phase of the simulation on the colony pair x, x' is completed without the disturbing effect of noise: even the zigging does not go beyond the boundary. Then the transfer phase begins which enters the disorder to the right of $x' + QB$.

We argue that there are only two ways for the head to get back to $x - c_{\text{attack}}QB$.

at1) The transfer into a new colony pair with starting point $y \geq x' + QB$ succeeds despite the disorder, and the clean interval extends over it, before the head moves left to $x - c_{\text{attack}}QB$ in the course of the regular simulation. Some inconsistencies may be discovered along the way, but they are corrected by healing.

at2) The inconsistencies encountered along the way trigger some rebuilding processes. Eventually, a complete, clean rebuilding area is created, the rebuilding succeeds, leaving a clean colony also to the right of $x' + QB$.

The Spill Bound property guarantees that the area to the left of $x' + (Q - c_{\text{spill}})B$ remains clean, therefore the only way for the head to move left of that is by the rules. Suppose that rebuilding is not initiated (with creating a substantial germ): then moving left can only happen by the normal course of simulation: the transfer stage of the simulation must be carried out, and this requires at least as many attacks to the right as the number of sweeps in the transfer stage. Every attack (followed by return) extends the clean interval further, until the whole target colony becomes clean, and the transfer completed. This is the case (at1).

Recall the rebuilding procedure in Section 9.3: the rebuilding area extends $2.5Q$ cells to the left and right from its initiating cell. This may become as large as $5QB$ to the left and right. If rebuilding is initiated, its starting position is necessarily to the right of $z = x' + (Q - \text{PadLen} - 2)B$. It then may extend to the left to at most $z - 5QB$ (this is overcounting, since the cells of the colony of x are all adjacent). Its many sweeps will result in attacks that clean an area to the right of the starting point. The procedure may be restarted several times, but those restartings will also be initiated to the right of z . Therefore the rebuilding area does not extend to the left of $z - 5QB$: if the head moves to the left of this, then the rebuilding must have succeeded. The rebuilding also must find or create a colony manifestly to the right of the restarting site: this will be to the right of z , moving this way the boundary of the area clean in M^* by at least QB : this is the case of (at2).

In the process described above, it is possible that the rebuilding finds a competing colony C starting at some $y \in x' + QB + [-\beta B, 0)$ which slightly (by the size of an island) overlaps from the right with the colony of x . The rebuilding

may decide to keep C and to overwrite the rest of the colony of x' as a bridge (or even target). This does not affect the result. \square

13 After a large burst

Our goal is to show that the simulation $M \rightarrow M^*$ defined in Section 10.3 is indeed a simulation. Section 11 shows this as long as the head operates in an area that is clean for the simulated machine M^* (can be super-annotated), and has no noise for machine M^* (that is its bursts on the level of machine M are isolated). In other words, essentially the Transition Function property of Definition 6.10 of trajectories for the simulated machine M^* has been taken care of.

The new element is the possibility of large areas that cannot be super-annotated: they may not even be clean, even on the level of machine M . The Spill Bound, Attack Cleaning, Dwell Cleaning and Pass Cleaning properties still must be proven.

13.1 Spill bound

One of the most complex analyses of this work is the proof that the simulated machine M^* also obeys the Spill Bound. Let us outline the problem.

We are looking at the boundary z of a large area that is clean for the machine M^* : without loss of generality suppose that this is the right boundary. We will be looking at it in a space-time rectangle $[a, b) \times [u, v)$ that is noise-free in M^* . The interesting case has $a < z < b$ with $z - a, b - z = O(QB)$. The assumption allows occasional bursts of noise of the trajectory η of M , but no two of these bursts must occur in a space-time rectangle of size comparable to the size of a colony work period of M . The Spill Bound for trajectory η keeps the disorder of η within $O(B)$ on the left of z while η is noise-free. If we could assume η noise-free then the heal/rebuild procedures would also keep the disorder of η^* from spilling over by more than $O(B^*) = O(QB)$. However, nothing is assumed about the length of the time interval $[u, v)$. If the head would spend all the time within $O(QB)$ of z then due to the Dwell Cleaning property of trajectories, the area would be cleaned out, again preventing spilling. But the head can slide out far to the right of b fast, since the disordered area on the right of z is arbitrarily large. Then it can come back much later to the left of z , and by a burst (allowed since much time has passed) can deposit an island of disorder there. Repeating this process would produce unlimited spillover, not only of the disorder of η^* but even that of η .

This is where the Pass Cleaning property helps. If the above process is repeated π times, the π passes would clean out an interval on the right of z whose size is of the order of QB , while depositing only π islands of disorder to the left of z . Our construction will have $\pi \ll Q$: more precisely, in our hierarchy of generalized Turing machines we will have $\pi_k = k$, $Q_k = k^2$. And $\ll Q$ bursts will still be handled by healing/rebuilding in η .

Of course this sketch is very crude, but it should help motivate the reasoning that follows.

Bibliography

- [1] Ilir Çapuni and Peter Gács. A Turing machine resisting isolated bursts of faults. *Chicago Journal of Theoretical Computer Science*, 2013. See also in arXiv:1203.1335. Extended abstract appeared in SOFSEM 2012. [2](#), [10.2](#)
- [2] Bruno Durand, Andrei E. Romashchenko, and Alexander Kh. Shen. Fixed-point tile sets and their applications. *Journal of Computer and System Sciences*, 78:731–764, 2012. [2.2](#)
- [3] Peter Gács. Reliable computation with cellular automata. *Journal of Computer System Science*, 32(1):15–78, February 1986. Conference version at STOC’ 83. [2.2](#)
- [4] Peter Gács. Reliable cellular automata with self-organization. *Journal of Statistical Physics*, 103(1/2):45–267, April 2001. See also arXiv:math/0003117 [math.PR] and the proceedings of STOC ’97. [1.4](#), [2.2](#)
- [5] G. L. Kurdyumov. An example of a nonergodic homogenous one-dimensional random medium with positive transition probabilities. *Soviet Mathematics Doklady*, 19(1):211–214, 1978. [2.2](#)